

Index

Symbols

- @ operator (matrix multiplication), [The Normal Equation](#)
- β (momentum), [Momentum](#)
- γ (gamma) value, [Gaussian RBF Kernel](#)
- ϵ (tolerance), [Batch Gradient Descent](#), [SVM Classes and Computational Complexity](#)
- ϵ greedy policy, [Exploration Policies](#)
- ϵ neighborhood, [DBSCAN](#)
- ϵ sensitive, [SVM Regression](#)
- χ^2 test, [Regularization Hyperparameters](#)
- ℓ_0 norm, [Select a Performance Measure](#)
- ℓ_1 norm, [Select a Performance Measure](#)
- ℓ_2 norm, [Select a Performance Measure](#)
- ℓ_k norm, [Select a Performance Measure](#)

A

- A/B experiments, [Training and Deploying TensorFlow Models at Scale](#)
- accelerated k-means, [Accelerated k-means and mini-batch k-means](#)
- accelerated linear algebra (XLA), [TensorFlow Functions and Graphs](#)
- accuracy performance measure, [What Is Machine Learning?](#),
[Measuring Accuracy Using Cross-Validation](#)
- ACF (autocorrelation function), [The ARMA Model Family](#)
- action advantage, reinforcement learning, [Evaluating Actions: The Credit Assignment Problem](#)
- action potentials (APs), [Biological Neurons](#)
- actions, in reinforcement learning, [Learning to Optimize Rewards](#),
[Evaluating Actions: The Credit Assignment Problem-Evaluating Actions: The Credit Assignment Problem](#)
- activation functions, [The Perceptron](#), [The Multilayer Perceptron and Backpropagation-The Multilayer Perceptron and Backpropagation](#)
 - for Conv2D layer, [Implementing Convolutional Layers with Keras](#)
 - custom models, [Custom Activation Functions, Initializers, Regularizers, and Constraints](#)

- ELU, [ELU and SELU-GELU, Swish, and Mish](#)
- GELU, [GELU, Swish, and Mish-GELU, Swish, and Mish](#)
- hyperbolic tangent (htan), [The Multilayer Perceptron and Backpropagation, Fighting the Unstable Gradients Problem](#)
- hyperparameters, [Learning Rate, Batch Size, and Other Hyperparameters](#)
- initialization parameters, [Glorot and He Initialization](#)
- leakyReLU, [Leaky ReLU-Leaky ReLU](#)
- Mish, [GELU, Swish, and Mish](#)
- PReLU, [Leaky ReLU](#)
- ReLU (see ReLU)
- RReLU, [Leaky ReLU](#)
- SELU, [ELU and SELU, Dropout](#)
- sigmoid, [Estimating Probabilities, The Multilayer Perceptron and Backpropagation, The Vanishing/Exploding Gradients Problems, Sparse Autoencoders](#)
- SiLU, [GELU, Swish, and Mish](#)
- softmax, [Softmax Regression, Classification MLPs, Creating the model using the sequential API, An Encoder–Decoder Network for Neural Machine Translation](#)
- softplus, [Regression MLPs](#)
- Swish, [GELU, Swish, and Mish](#)
- active learning, [Using Clustering for Semi-Supervised Learning](#)
- actor-critic, [Overview of Some Popular RL Algorithms](#)
- actual class, [Confusion Matrices](#)
- actual versus estimated probabilities, [The ROC Curve](#)
- AdaBoost, [AdaBoost-AdaBoost](#)
- AdaGrad, [AdaGrad](#)
- Adam optimization, [Adam, \$\ell_1\$ and \$\ell_2\$ Regularization](#)
- AdaMax, [AdaMax](#)
- AdamW, [AdamW, \$\ell_1\$ and \$\ell_2\$ Regularization](#)
- adaptive boosting (AdaBoost), [Boosting-AdaBoost](#)
- adaptive instance normalization (AdaIN), [StyleGANs](#)
- adaptive learning rate algorithms, [AdaGrad-AdamW](#)
- adaptive moment estimation (Adam), [Adam](#)
- additive attention, [Attention Mechanisms](#)
- advantage actor-critic (A2C), [Overview of Some Popular RL Algorithms](#)
- adversarial learning, [Semantic Segmentation, Autoencoders, GANs, and Diffusion Models](#)

- affine transformations, [StyleGANs](#)
- affinity function, [Clustering Algorithms: k-means and DBSCAN](#)
- affinity propagation, [Other Clustering Algorithms](#)
- agents, reinforcement learning, [Reinforcement learning](#), [Learning to Optimize Rewards](#), [Policy Gradients](#), [Q-Learning](#)
- agglomerative clustering, [Other Clustering Algorithms](#)
- Akaike information criterion (AIC), [Selecting the Number of Clusters](#)
- AlexNet, [AlexNet](#)
- algorithms, preparing data for, [Prepare the Data for Machine Learning Algorithms-Transformation Pipelines](#)
- alignment model, [Attention Mechanisms](#)
- AllReduce algorithm, [Data parallelism using the mirrored strategy](#)
- alpha dropout, [Dropout](#)
- AlphaGo, [Reinforcement learning](#), [Reinforcement Learning](#), [Overview of Some Popular RL Algorithms](#)
- anchor priors, [You Only Look Once](#)
- ANNs (see artificial neural networks)
- anomaly detection, [Examples of Applications](#), [Unsupervised learning](#)
 - clustering for, [Clustering Algorithms: k-means and DBSCAN](#)
 - GMM, [Using Gaussian Mixtures for Anomaly Detection-Using Gaussian Mixtures for Anomaly Detection](#)
 - isolation forest, [Other Algorithms for Anomaly and Novelty Detection](#)
- AP (average precision), [You Only Look Once](#)
- APs (action potentials), [Biological Neurons](#)
- area under the curve (AUC), [The ROC Curve](#)
- argmax(), [Softmax Regression](#)
- ARIMA model, [The ARMA Model Family-The ARMA Model Family](#)
- ARMA model family, [The ARMA Model Family-The ARMA Model Family](#)
- artificial neural networks (ANNs), [Introduction to Artificial Neural Networks with Keras-Learning Rate, Batch Size, and Other Hyperparameters](#)
 - autoencoders (see autoencoders)
 - backpropagation, [The Multilayer Perceptron and Backpropagation-The Multilayer Perceptron and Backpropagation](#)
 - biological neurons as background to, [Biological Neurons-Biological Neurons](#)

- evolution of, [From Biological to Artificial Neurons-From Biological to Artificial Neurons](#)
- hyperparameter fine-tuning, [Fine-Tuning Neural Network Hyperparameters-Learning Rate, Batch Size, and Other Hyperparameters](#)
- implementing MLPs with Keras, [Implementing MLPs with Keras-Using TensorBoard for Visualization](#)
- logical computations with neurons, [Logical Computations with Neurons-Logical Computations with Neurons](#)
- perceptrons (see multilayer perceptrons)
- reinforcement learning policies, [Neural Network Policies](#)
- artificial neuron, [Logical Computations with Neurons](#)
- association rule learning, [Unsupervised learning](#)
- assumptions, checking in model building, [Data Mismatch, Check the Assumptions](#)
- asynchronous advantage actor-critic (A3C), [Overview of Some Popular RL Algorithms](#)
- asynchronous gang scheduling, [Bandwidth saturation](#)
- asynchronous updates, with centralized parameters, [Asynchronous updates](#)
- à-trous convolutional layer, [Semantic Segmentation](#)
- attention mechanisms, [Natural Language Processing with RNNs and Attention, Attention Mechanisms-Multi-head attention, Vision Transformers](#)
 - (see also transformer models)
- attributes, [Take a Quick Look at the Data Structure-Take a Quick Look at the Data Structure](#)
 - categorical, [Handling Text and Categorical Attributes, Handling Text and Categorical Attributes](#)
 - combinations of, [Experiment with Attribute Combinations-Experiment with Attribute Combinations](#)
 - preprocessed, [Take a Quick Look at the Data Structure](#)
 - target, [Take a Quick Look at the Data Structure](#)
 - unsupervised learning, [Supervised learning](#)
- AUC (area under the curve), [The ROC Curve](#)
- autocorrelated time series, [Forecasting a Time Series](#)
- autocorrelation function (ACF), [The ARMA Model Family](#)
- autodiff (automatic differentiation), [Autodiff-Reverse-Mode Autodiff](#)

- for computing gradients, [Computing Gradients Using Autodiff](#)-
[Computing Gradients Using Autodiff](#)
- finite difference approximation, [Finite Difference Approximation](#)
- forward-mode, [Forward-Mode Autodiff](#)-[Forward-Mode Autodiff](#)
- manual differentiation, [Manual Differentiation](#)
- reverse-mode, [Reverse-Mode Autodiff](#)-[Reverse-Mode Autodiff](#)
- autoencoders, [Autoencoders, GANs, and Diffusion Models](#)-[Generating Fashion MNIST Images](#)
 - convolutional, [Convolutional Autoencoders](#)-[Convolutional Autoencoders](#)
 - denoising, [Denoising Autoencoders](#)-[Denoising Autoencoders](#)
 - efficient data representations, [Efficient Data Representations](#)-[Efficient Data Representations](#)
 - overcomplete, [Convolutional Autoencoders](#)
 - PCA with undercomplete linear autoencoder, [Performing PCA with an Undercomplete Linear Autoencoder](#)-[Performing PCA with an Undercomplete Linear Autoencoder](#)
 - sparse, [Sparse Autoencoders](#)-[Sparse Autoencoders](#)
 - stacked, [Stacked Autoencoders](#)-[Training One Autoencoder at a Time](#)
 - training one at a time, [Training One Autoencoder at a Time](#)-[Training One Autoencoder at a Time](#)
 - undercomplete, [Efficient Data Representations](#)
 - variational, [Variational Autoencoders](#)-[Variational Autoencoders](#)
- Autograph, [AutoGraph and Tracing](#)
- AutoGraph, [Using AutoGraph to Capture Control Flow](#)
- AutoML service, [Fine-Tuning Neural Network Hyperparameters](#),
[Hyperparameter Tuning on Vertex AI](#)
- autoregressive integrated moving average (ARIMA) model, [The ARMA Model Family](#)-[The ARMA Model Family](#)
- autoregressive model, [The ARMA Model Family](#)
- autoregressive moving average (ARMA) model family, [The ARMA Model Family](#)-[The ARMA Model Family](#)
- auxiliary task, pretraining on, [Pretraining on an Auxiliary Task](#)
- average absolute deviation, [Select a Performance Measure](#)
- average pooling layer, [Implementing Pooling Layers with Keras](#)
- average precision (AP), [You Only Look Once](#)

B

- backbone, model, [GoogLeNet](#)

- backpropagation, [The Multilayer Perceptron and Backpropagation](#)-
[The Multilayer Perceptron and Backpropagation](#), [The Vanishing/Exploding Gradients Problems](#), [Computing Gradients Using Autodiff](#), [Generative Adversarial Networks](#)
- backpropagation through time (BPTT), [Training RNNs](#)
- bagging (bootstrap aggregating), [Bagging and Pasting](#)-[Random Patches and Random Subspaces](#)
- Bahdanau attention, [Attention Mechanisms](#)
- balanced iterative reducing and clustering using hierarchies (BIRCH),
[Other Clustering Algorithms](#)
- bandwidth saturation, [Bandwidth saturation](#)-[Bandwidth saturation](#)
- BaseEstimator, [Custom Transformers](#)
- batch gradient descent, [Batch Gradient Descent](#)-[Batch Gradient Descent](#), [Ridge Regression](#)
- batch learning, [Batch learning](#)-[Batch learning](#)
- batch normalization (BN), [Batch Normalization](#)-[Implementing batch normalization with Keras](#), [Fighting the Unstable Gradients Problem](#)
- batch predictions, [Deploying a new model version](#), [Creating a Prediction Service on Vertex AI](#), [Running Batch Prediction Jobs on Vertex AI](#)
- Bayesian Gaussian mixtures, [Bayesian Gaussian Mixture Models](#)
- Bayesian information criterion (BIC), [Selecting the Number of Clusters](#)
- beam search, [Beam Search](#)-[Beam Search](#)
- Bellman optimality equation, [Markov Decision Processes](#)
- bias, [Learning Curves](#)
- bias and fairness, NLP transformers, [Hugging Face's Transformers Library](#)
- bias term constant, [Linear Regression](#), [The Normal Equation](#)
- bias/variance trade-off, [Learning Curves](#)
- BIC (Bayesian information criterion), [Selecting the Number of Clusters](#)
- bidirectional recurrent layer, [Bidirectional RNNs](#)
- binary classifiers, [Training a Binary Classifier](#), [Logistic Regression](#)
- binary logarithm, [Computational Complexity](#)
- binary trees, [Making Predictions](#), [Other Clustering Algorithms](#)
- biological neural networks (BNNs), [Biological Neurons](#)-[Biological Neurons](#)
- BIRCH (balanced iterative reducing and clustering using hierarchies),
[Other Clustering Algorithms](#)
- black box models, [Making Predictions](#)

- blending, in stacking, [Stacking](#)
- BN (batch normalization), [Batch Normalization](#), [Fighting the Unstable Gradients Problem](#)
- BNNs (biological neural networks), [Biological Neurons-Biological Neurons](#)
- boosting, [Boosting-Histogram-Based Gradient Boosting](#)
- bootstrap aggregating (bagging), [Bagging and Pasting-Random Patches and Random Subspaces](#)
- bootstrapping, [Bagging and Pasting](#)
- bottleneck layers, [GoogLeNet](#)
- bounding boxes, image identification, [Classification and Localization-You Only Look Once](#)
- BPTT (backpropagation through time), [Training RNNs](#)
- bucketizing a feature, [Feature Scaling and Transformation](#)
- byte pair encoding (BPE), [Sentiment Analysis](#)

C

- California Housing Prices dataset, [Working with Real Data-Check the Assumptions](#)
- callbacks, [Using Callbacks-Using Callbacks](#)
- CART (Classification and Regression Tree) algorithm, [Making Predictions](#), [The CART Training Algorithm](#), [Regression](#)
- catastrophic forgetting, [Implementing Deep Q-Learning](#)
- categorical attributes, [Handling Text and Categorical Attributes](#), [Handling Text and Categorical Attributes](#)
- categorical features, encoding, [The Hashing Layer-Encoding Categorical Features Using Embeddings](#)
- CategoryEncoding layer, [The CategoryEncoding Layer](#)
- causal model, [Forecasting Using a Sequence-to-Sequence Model](#), [Bidirectional RNNs](#)
- centralized parameters, [Data parallelism with centralized parameters-Asynchronous updates](#)
- centroid, cluster, [Clustering Algorithms: k-means and DBSCAN](#), [k-means](#), [The k-means algorithm-Centroid initialization methods](#)
- chain rule, [The Multilayer Perceptron and Backpropagation](#)
- chain-of-thought prompting, [An Avalanche of Transformer Models](#)
- ChainClassifier, [Multilabel Classification](#)
- chaining transformations, [Chaining Transformations-Chaining Transformations](#)

- char-RNN model, [Generating Shakespearean Text Using a Character RNN-Stateful RNN](#)
- chatbot or personal assistant, [Examples of Applications](#)
- check_estimator(), [Custom Transformers](#)
- chi-squared (χ^2) test, [Regularization Hyperparameters](#)
- classification, [Classification-Multioutput Classification](#)
 - application examples of, [Examples of Applications-Examples of Applications](#)
 - binary classifier, [Training a Binary Classifier](#), [Logistic Regression](#)
 - CNNs, [Classification and Localization-You Only Look Once](#)
 - error analysis, [Error Analysis-Error Analysis](#)
 - hard margin, [Soft Margin Classification](#), [Under the Hood of Linear SVM Classifiers](#)
 - hard voting classifiers, [Voting Classifiers](#)
 - image (see images)
 - logistic regression (see logistic regression)
 - MLPs for, [Classification MLPs-Classification MLPs](#)
 - MNIST dataset, [MNIST-MNIST](#)
 - multiclass, [Multiclass Classification-Multiclass Classification](#), [Classification MLPs-Classification MLPs](#)
 - multilabel, [Multilabel Classification-Multilabel Classification](#)
 - multioutput, [Multioutput Classification-Multioutput Classification](#)
 - performance measures, [Performance Measures-The ROC Curve](#)
 - and regression, [Supervised learning](#), [Multioutput Classification](#)
 - soft margin, [Soft Margin Classification-Soft Margin Classification](#)
 - softmax regression, [Softmax Regression-Softmax Regression](#)
 - SVMs (see support vector machines)
 - text, [Sentiment Analysis-Reusing Pretrained Embeddings and Language Models](#)
 - voting classifiers, [Voting Classifiers-Voting Classifiers](#)
- Classification and Regression Tree (CART) algorithm, [Making Predictions](#), [The CART Training Algorithm](#), [Regression](#)
- clone(), [Early Stopping](#)
- closed-form equation/solution, [Training Models](#), [The Normal Equation](#), [Ridge Regression](#), [Training and Cost Function](#)
- cloud platform deployment with Vertex AI, [Creating a Prediction Service on Vertex AI](#)-[Creating a Prediction Service on Vertex AI](#)
- clustering algorithms, [Examples of Applications](#), [Unsupervised learning](#), [Unsupervised Learning Techniques-Other Clustering Algorithms](#)

- affinity propagation, [Other Clustering Algorithms](#)
- agglomerative clustering, [Other Clustering Algorithms](#)
- applications for, [Clustering Algorithms: k-means and DBSCAN](#)-
[Clustering Algorithms: k-means and DBSCAN](#)
- BIRCH, [Other Clustering Algorithms](#)
- DBSCAN, [DBSCAN](#)
- GMM, [Gaussian Mixtures](#)-[Other Algorithms for Anomaly and Novelty Detection](#)
- image segmentation, [Clustering Algorithms: k-means and DBSCAN](#),
[Using Clustering for Image Segmentation](#)-[Using Clustering for Image Segmentation](#)
- k-means (see k-means algorithm)
- mean-shift, [Other Clustering Algorithms](#)
- responsibilities of clusters for instances, [Gaussian Mixtures](#)
- semi-supervised learning with, [Using Clustering for Semi-Supervised Learning](#)-[Using Clustering for Semi-Supervised Learning](#)
- spectral clustering, [Other Clustering Algorithms](#)
- CNNs (see convolutional neural networks)
- Colab, [Running the Code Examples Using Google Colab](#)-[Running the Code Examples Using Google Colab](#)
- color channels, [Stacking Multiple Feature Maps](#)
- color segmentation, images, [Using Clustering for Image Segmentation](#)
- column vectors, [Linear Regression](#)
- ColumnTransformer, [Transformation Pipelines](#)
- complex models with functional API, [Building Complex Models Using the Functional API](#)-[Building Complex Models Using the Functional API](#)
- compound scaling, [Other Noteworthy Architectures](#)
- compressed TFRecord files, [Compressed TFRecord Files](#)
- compression and decompression, PCA, [PCA for Compression](#)-[PCA for Compression](#)
- computation graphs, [A Quick Tour of TensorFlow](#)
- computational complexity
 - DBSCAN, [DBSCAN](#)
 - decision trees, [Computational Complexity](#)
 - Gaussian mixture model, [Gaussian Mixtures](#)
 - histogram-based gradient boosting, [Histogram-Based Gradient Boosting](#)
 - k-means algorithm, [The k-means algorithm](#)

- Normal equation, [Computational Complexity](#)
- and SVM classes, [SVM Classes and Computational Complexity](#)
- concatenative attention, [Attention Mechanisms](#)
- concrete function, [TF Functions and Concrete Functions-Exploring Function Definitions and Graphs](#)
- conditional GAN, [Deep Convolutional GANs](#)
- conditional probability, [Beam Search](#)
- confidence interval, [Evaluate Your System on the Test Set](#)
- confusion matrix (CM), [Confusion Matrices-Confusion Matrices, Error Analysis-Error Analysis](#)
- ConfusionMatrixDisplay, [Error Analysis](#)
- connectionism, [From Biological to Artificial Neurons](#)
- constrained optimization, [Under the Hood of Linear SVM Classifiers](#)
- constraints, custom models, [Custom Activation Functions, Initializers, Regularizers, and Constraints](#)
- convergence rate, [Batch Gradient Descent](#)
- convex function, [Gradient Descent](#)
- convex quadratic optimization, [Under the Hood of Linear SVM Classifiers](#)
- convolution kernels (kernels), [Filters, CNN Architectures](#)
- convolutional neural networks (CNNs), [Deep Computer Vision Using Convolutional Neural Networks-Semantic Segmentation](#)
 - architectures, [CNN Architectures-Choosing the Right CNN Architecture](#)
 - autoencoders, [Convolutional Autoencoders-Convolutional Autoencoders](#)
 - classification and localization, [Classification and Localization-You Only Look Once](#)
 - convolutional layers, [Convolutional Layers-Memory Requirements, Semantic Segmentation-Semantic Segmentation, Using 1D convolutional layers to process sequences, WaveNet-WaveNet, Masking](#)
 - evolution of, [Deep Computer Vision Using Convolutional Neural Networks](#)
 - GANs, [Deep Convolutional GANs-Deep Convolutional GANs](#)
 - object detection, [Object Detection-You Only Look Once](#)
 - object tracking, [Object Tracking](#)
 - pooling layers, [Pooling Layers-Implementing Pooling Layers with Keras](#)

- pretrained models from Keras, [Using Pretrained Models from Keras](#)
[Using Pretrained Models from Keras](#)
- ResNet-34 CNN using Keras, [Implementing a ResNet-34 CNN Using Keras](#)
- semantic segmentation, [Examples of Applications](#), [Using Clustering for Image Segmentation](#), [Semantic Segmentation-Semantic Segmentation](#)
- splitting across devices, [Model Parallelism](#)
- transfer learning pretrained models, [Pretrained Models for Transfer Learning](#)-[Pretrained Models for Transfer Learning](#)
- U-Net, [Diffusion Models](#)
- and vision transformers, [Vision Transformers](#)
- visual cortex architecture, [The Architecture of the Visual Cortex](#)
- WaveNet, [WaveNet-WaveNet](#)
- copy.deepcopy(), [Early Stopping](#)
- core instance, [DBSCAN](#)
- correlation coefficient, [Look for Correlations](#)-[Look for Correlations](#)
- cost function, [Model-based learning and a typical machine learning workflow](#), [Select a Performance Measure](#)
 - in AdaBoost, [AdaBoost](#)
 - in autoencoders, [Performing PCA with an Undercomplete Linear Autoencoder](#)
 - in bounding box prediction model, [Classification and Localization](#)
 - in CART training algorithm, [The CART Training Algorithm](#), [Regression](#)
 - in elastic net, [Elastic Net Regression](#)
 - in gradient descent, [Training Models](#), [Gradient Descent-Gradient Descent](#), [Stochastic Gradient Descent](#)-[Stochastic Gradient Descent](#), [The Vanishing/Exploding Gradients Problems](#)
 - in lasso regression, [Lasso Regression](#)-[Lasso Regression](#)
 - in Nesterov accelerated gradient, [Nesterov Accelerated Gradient](#)
 - in linear regression, [Linear Regression](#)
 - in logistic regression, [Training and Cost Function](#)-[Training and Cost Function](#)
 - in momentum optimization, [Momentum](#)
 - in ridge regression, [Ridge Regression](#)
 - in variational autoencoders, [Variational Autoencoders](#)
- credit assignment problem, [Evaluating Actions: The Credit Assignment Problem](#)-[Evaluating Actions: The Credit Assignment Problem](#)

- cross entropy, [Softmax Regression](#)
- cross-validation, [Hyperparameter Tuning and Model Selection](#), [Better Evaluation Using Cross-Validation](#)-[Better Evaluation Using Cross-Validation](#), [Evaluate Your System on the Test Set](#), [Measuring Accuracy Using Cross-Validation](#)-[Measuring Accuracy Using Cross-Validation](#), [Learning Curves](#)-[Learning Curves](#)
- cross_val_predict(), [Confusion Matrices](#), [The Precision/Recall Trade-off](#), [The ROC Curve](#), [Error Analysis](#), [Stacking](#)
- cross_val_score(), [Better Evaluation Using Cross-Validation](#), [Measuring Accuracy Using Cross-Validation](#)
- CUDA library, [Getting Your Own GPU](#)
- curiosity-based exploration, [Overview of Some Popular RL Algorithms](#)
- curriculum learning, [Overview of Some Popular RL Algorithms](#)
- custom models and training algorithms, [Customizing Models and Training Algorithms](#)-[Custom Training Loops](#)
 - activation functions, [Custom Activation Functions](#), [Initializers](#), [Regularizers, and Constraints](#)
 - autodiff for computing gradients, [Computing Gradients Using Autodiff](#)-[Computing Gradients Using Autodiff](#)
 - constraints, [Custom Activation Functions](#), [Initializers](#), [Regularizers, and Constraints](#)
 - initializers, [Custom Activation Functions](#), [Initializers](#), [Regularizers, and Constraints](#)
 - layers, [Custom Layers](#)-[Custom Layers](#)
 - loss functions, [Custom Loss Functions](#), [Losses and Metrics Based on Model Internals](#)-[Losses and Metrics Based on Model Internals](#)
 - metrics, [Custom Metrics](#)-[Custom Metrics](#), [Losses and Metrics Based on Model Internals](#)-[Losses and Metrics Based on Model Internals](#)
 - models, [Custom Models](#)-[Custom Models](#)
 - regularizers, [Custom Activation Functions](#), [Initializers](#), [Regularizers, and Constraints](#)
 - saving and loading models, [Saving and Loading Models That Contain Custom Components](#)-[Saving and Loading Models That Contain Custom Components](#)
 - training loops, [Custom Training Loops](#)-[Custom Training Loops](#)
- custom transformers, [Custom Transformers](#)-[Custom Transformers](#)
- customer segmentation, [Clustering Algorithms: k-means and DBSCAN](#)

- DALL-E, [Vision Transformers](#)
- data
 - downloading, [Saving Your Code Changes and Your Data](#)-[Saving Your Code Changes and Your Data](#)
 - efficient data representations, [Efficient Data Representations](#)-[Efficient Data Representations](#)
 - enqueueing and dequeuing, [Queues](#)
 - finding correlations in, [Look for Correlations](#)-[Look for Correlations](#)
 - making assumptions about, [Data Mismatch](#)
 - overfitting (see overfitting of data)
 - preparing for ML algorithms, [Prepare the Data for Machine Learning Algorithms](#)-[Transformation Pipelines](#)
 - preparing for ML models, [Preparing the Data for Machine Learning Models](#)-[Preparing the Data for Machine Learning Models](#)
 - preprocessing (see loading and preprocessing data)
 - shuffling of, [Shuffling the Data](#)
 - test data (see test set)
 - time series (see time series data)
 - training data (see training set)
 - underfitting of, [Train and Evaluate on the Training Set](#), [Learning Curves](#)-[Learning Curves](#), [Polynomial Kernel](#)
 - unreasonable effectiveness, [Insufficient Quantity of Training Data](#)
 - visualizing (see visualization)
 - working with real data, [Working with Real Data](#)-[Working with Real Data](#)
- data analysis, clustering for, [Clustering Algorithms](#): k-means and [DBSCAN](#)
- data augmentation, [Exercises](#), [AlexNet](#), [Pretrained Models for Transfer Learning](#)
- data cleaning, [Clean the Data](#)-[Clean the Data](#)
- data drift, [Batch learning](#)
- data mining, [Why Use Machine Learning?](#)
- data mismatch, [Data Mismatch](#)
- data parallelism, [Data Parallelism](#)-[Bandwidth saturation](#)
- data pipeline, [Frame the Problem](#)
- data snooping bias, [Create a Test Set](#)
- data structure, [Take a Quick Look at the Data Structure](#)-[Take a Quick Look at the Data Structure](#), [Special Data Structures](#)-[Queues](#)
- data-efficient image transformers (DeiT), [Vision Transformers](#)

- DataFrame, [Clean the Data](#), [Handling Text and Categorical Attributes](#), [Feature Scaling and Transformation](#)
- Dataquest, [Other Resources](#)
- Datasets library, [The TensorFlow Datasets Project](#)
- DBSCAN, [DBSCAN-DBSCAN](#)
- DCGANs (deep convolutional GANS), [Deep Convolutional GANs-Deep Convolutional GANs](#)
- DDPM (denoising diffusion probabilistic model), [Autoencoders, GANs, and Diffusion Models](#), [Diffusion Models-Diffusion Models](#)
- DDQN (dueling DQN), [Dueling DQN](#)
- decision boundaries, [Decision Boundaries-Decision Boundaries](#), [Softmax Regression](#), [Making Predictions](#), [Sensitivity to Axis Orientation](#), [k-means](#)
- decision function, [The Precision/Recall Trade-off](#), [Under the Hood of Linear SVM Classifiers-Under the Hood of Linear SVM Classifiers](#)
- decision stumps, [AdaBoost](#)
- decision threshold, [The Precision/Recall Trade-off](#)-[The Precision/Recall Trade-off](#)
- decision trees, [Decision Trees-Decision Trees Have a High Variance](#), [Ensemble Learning and Random Forests](#)
 - (see also ensemble learning)
 - bagging and pasting, [Bagging and Pasting in Scikit-Learn](#)
 - CART training algorithm, [The CART Training Algorithm](#)
 - class probability estimates, [Estimating Class Probabilities](#)
 - computational complexity, [Computational Complexity](#)
 - and decision boundaries, [Making Predictions](#)
 - GINI impurity or entropy measures, [Gini Impurity or Entropy?](#)
 - high variance with, [Decision Trees Have a High Variance](#)
 - predictions, [Making Predictions-The CART Training Algorithm](#)
 - regression tasks, [Regression-Regression](#)
 - regularization hyperparameters, [Regularization Hyperparameters-Regularization Hyperparameters](#)
 - sensitivity to axis orientation, [Sensitivity to Axis Orientation](#)
 - training and visualizing, [Training and Visualizing a Decision Tree](#)-[Training and Visualizing a Decision Tree](#)
 - in training the model, [Train and Evaluate on the Training Set-Better Evaluation Using Cross-Validation](#)
- DecisionTreeClassifier, [Training and Visualizing a Decision Tree](#), [Gini Impurity or Entropy?](#), [Regularization Hyperparameters](#), [Sensitivity to](#)

Axis Orientation, Random Forests

- DecisionTreeRegressor, [Train and Evaluate on the Training Set](#), [Decision Trees](#), [Regression](#), [Gradient Boosting](#)
- decision_function(), [The Precision/Recall Trade-off](#)
- deconvolution layer, [Semantic Segmentation](#)
- deep autoencoders (see stacked autoencoders)
- deep convolutional GANS (DCGANs), [Deep Convolutional GANs](#)-[Deep Convolutional GANs](#)
- deep Gaussian process, [Monte Carlo \(MC\) Dropout](#)
- deep learning, [The Machine Learning Tsunami](#)
 - (see also deep neural networks; reinforcement learning)
- deep neural networks (DNNs), [Training Deep Neural Networks](#)-[Summary and Practical Guidelines](#)
 - CNNs (see convolutional neural networks)
 - default configuration, [Summary and Practical Guidelines](#)
 - faster optimizers for, [Faster Optimizers](#)-[AdamW](#)
 - learning rate scheduling, [Learning Rate Scheduling](#)-[Learning Rate Scheduling](#)
 - MLPs (see multilayer perceptrons)
 - regularization, [Avoiding Overfitting Through Regularization](#)-[Max-Norm Regularization](#)
 - reusing pretrained layers, [Reusing Pretrained Layers](#)-[Pretraining on an Auxiliary Task](#)
 - RNNs (see recurrent neural networks)
 - and transfer learning, [Self-supervised learning](#)
 - unstable gradients, [The Vanishing/Exploding Gradients Problems](#)
 - vanishing and exploding gradients, [The Vanishing/Exploding Gradients Problems](#)-[Gradient Clipping](#)
- deep neuroevolution, [Fine-Tuning Neural Network Hyperparameters](#)
- deep Q-learning, [Approximate Q-Learning](#) and [Deep Q-Learning](#)-[Dueling DQN](#)
- deep Q-networks (DQNs) (see Q-learning algorithm)
- deepcopy(), [Early Stopping](#)
- DeepMind, [Reinforcement learning](#)
- degrees of freedom, [Overfitting the Training Data](#), [Learning Curves](#)
- DeiT (data-efficient image transformers), [Vision Transformers](#)
- denoising autoencoders, [Denoising Autoencoders](#)-[Denoising Autoencoders](#)

- denoising diffusion probabilistic model (DDPM), [Autoencoders, GANs, and Diffusion Models](#), [Diffusion Models-Diffusion Models](#)
- Dense layer, [The Perceptron](#), [Creating the model using the sequential API](#), [Creating the model using the sequential API](#), [Building Complex Models Using the Functional API](#)
- dense matrix, [Feature Scaling and Transformation](#), [Transformation Pipelines](#)
- density estimation, [Unsupervised Learning Techniques](#), [DBSCAN-DBSCAN](#), [Gaussian Mixtures](#)
- density threshold, [Using Gaussian Mixtures for Anomaly Detection](#)
- depth concatenation layer, [GoogLeNet](#)
- depthwise separable convolution layer, [Xception](#)
- deque, [Implementing Deep Q-Learning](#)
- describe(), [Take a Quick Look at the Data Structure](#)
- development set (dev set), [Hyperparameter Tuning and Model Selection](#)
- differencing, time series forecasting, [Forecasting a Time Series](#), [Forecasting a Time Series](#), [The ARMA Model Family](#)
- diffusion models, [Autoencoders, GANs, and Diffusion Models](#), [Diffusion Models-Diffusion Models](#)
- dilated filter, [Semantic Segmentation](#)
- dilation rate, [Semantic Segmentation](#)
- dimensionality reduction, [Unsupervised learning](#), [Dimensionality Reduction-Other Dimensionality Reduction Techniques](#)
 - approaches to, [Main Approaches for Dimensionality Reduction-Manifold Learning](#)
 - autoencoders, [Autoencoders, GANs, and Diffusion Models](#), [Efficient Data Representations-Performing PCA with an Undercomplete Linear Autoencoder](#)
 - choosing the right number of dimensions, [Choosing the Right Number of Dimensions](#)
 - clustering, [Clustering Algorithms: k-means and DBSCAN](#)
 - curse of dimensionality, [The Curse of Dimensionality-The Curse of Dimensionality](#)
 - for data visualization, [Dimensionality Reduction](#)
 - information loss from, [Dimensionality Reduction](#)
 - Isomap, [Other Dimensionality Reduction Techniques](#)
 - linear discriminant analysis, [Other Dimensionality Reduction Techniques](#)

- LLE, [LLE-LLE](#)
- multidimensional scaling, [Other Dimensionality Reduction Techniques](#)
- PCA (see principal component analysis)
- random projection algorithm, [Random Projection-Random Projection](#)
- t-distributed stochastic neighbor embedding, [Other Dimensionality Reduction Techniques](#), [Visualizing the Fashion MNIST Dataset](#)
- discount factor γ in reward system, [Evaluating Actions: The Credit Assignment Problem](#)
- discounted rewards, [Evaluating Actions: The Credit Assignment Problem](#), [Policy Gradients](#)
- Discretization layer, [The Discretization Layer](#)
- discriminator, GAN, [Autoencoders, GANs, and Diffusion Models](#), [Generative Adversarial Networks-The Difficulties of Training GANs](#)
- DistilBERT model, [An Avalanche of Transformer Models](#), [Hugging Face's Transformers Library-Hugging Face's Transformers Library](#)
- distillation, [An Avalanche of Transformer Models](#)
- distribution strategies API, [Training at Scale Using the Distribution Strategies API](#)
- Docker container, [Installing and starting TensorFlow Serving](#)
- dot product, [Attention Mechanisms](#)
- Dota 2, [Overview of Some Popular RL Algorithms](#)
- Double DQN, [Double DQN](#)
- downloading data, [Saving Your Code Changes and Your Data-Saving Your Code Changes and Your Data](#)
- DQNs (deep Q-networks) (see Q-learning algorithm)
- drop(), [Prepare the Data for Machine Learning Algorithms](#)
- dropna(), [Clean the Data](#)
- Dropout, [Denoising Autoencoders](#)
- dropout rate, [Dropout](#)
- dropout regularization, [Dropout-Dropout](#)
- dual numbers, [Forward-Mode Autodiff](#)
- dual problem, [The Dual Problem-Kernelized SVMs](#)
- dueling DQN (DDQN), [Dueling DQN](#)
- dummy attributes, [Handling Text and Categorical Attributes](#)
- dying ReLU problem, [Better Activation Functions](#)
- dynamic models with subclassing API, [Using the Subclassing API to Build Dynamic Models-Using the Subclassing API to Build Dynamic](#)

Models

- dynamic programming, [Markov Decision Processes](#)

E

- eager execution (eager mode), [AutoGraph and Tracing](#)
- early stopping regularization, [Early Stopping-Early Stopping](#), [Gradient Boosting](#), [Forecasting Using a Linear Model](#)
- edge computing, [Deploying a Model to a Mobile or Embedded Device](#)
- efficient data representations, autoencoders, [Efficient Data Representations-Efficient Data Representations](#)
- elastic net, [Elastic Net Regression](#)
- EllipticEnvelope, [Other Algorithms for Anomaly and Novelty Detection](#)
- ELMo (Embeddings from Language Models), [Reusing Pretrained Embeddings and Language Models](#)
- ELU (exponential linear unit), [ELU and SELU-GELU, Swish, and Mish](#)
- EM (expectation-maximization), [Gaussian Mixtures](#)
- embedded device, deploying model to, [Deploying a Model to a Mobile or Embedded Device](#)-[Deploying a Model to a Mobile or Embedded Device](#)
- embedded Reber grammars, [Exercises](#)
- embedding matrix, [Encoding Categorical Features Using Embeddings](#), [An Encoder–Decoder Network for Neural Machine Translation](#)
- embedding size, [An Encoder–Decoder Network for Neural Machine Translation](#), [Positional encodings](#)
- embeddings, [Handling Text and Categorical Attributes](#)
 - encoding categorical features using, [Encoding Categorical Features Using Embeddings-Encoding Categorical Features Using Embeddings](#)
 - reusing pretrained, [Reusing Pretrained Embeddings and Language Models-Reusing Pretrained Embeddings and Language Models](#)
 - sentiment analysis, [Sentiment Analysis](#)
- Embeddings from Language Models (ELMo), [Reusing Pretrained Embeddings and Language Models](#)
- encoder, [Efficient Data Representations](#)
 - (see also autoencoders)
- encoder–decoder models, [Input and Output Sequences](#), [Natural Language Processing with RNNs and Attention](#), [An Encoder–Decoder Network for Neural Machine Translation](#)-[Beam Search](#)
 - (see also attention mechanisms)

- end-to-end ML project exercise, [End-to-End Machine Learning Project-Try It Out!](#)
 - building the model, [Look at the Big Picture-Check the Assumptions](#)
 - discovering and visualizing data, [Explore and Visualize the Data to Gain Insights-Experiment with Attribute Combinations](#)
 - fine-tuning your model, [Fine-Tune Your Model-Evaluate Your System on the Test Set](#)
 - getting the data, [Get the Data-Create a Test Set](#)
 - preparing data for ML algorithms, [Prepare the Data for Machine Learning Algorithms-Transformation Pipelines](#)
 - real data, advantages of working with, [Working with Real Data-Working with Real Data](#)
 - selecting and training a model, [Train and Evaluate on the Training Set-Train and Evaluate on the Training Set](#)
- endpoint, deploying model on GCP, [Creating a Prediction Service on Vertex AI](#)
- ensemble learning, [Ensemble Learning and Random Forests-Stacking](#)
 - bagging and pasting, [Bagging and Pasting-Random Patches and Random Subspaces](#)
 - boosting, [Boosting-Histogram-Based Gradient Boosting](#)
 - cross-validation, [Better Evaluation Using Cross-Validation](#)
 - fine-tuning the system, [Ensemble Methods](#)
 - random forests (see random forests)
 - stacking, [Stacking-Stacking](#)
 - voting classifiers, [Voting Classifiers-Voting Classifiers](#)
- entailment, [An Avalanche of Transformer Models](#)
- entropy impurity measure, [Gini Impurity or Entropy?](#)
- environments, reinforcement learning, [Learning to Optimize Rewards, Introduction to OpenAI Gym-Introduction to OpenAI Gym](#)
- epochs, [Batch Gradient Descent](#)
- equalized learning rate, GAN, [Progressive Growing of GANs](#)
- equivariance, [Pooling Layers](#)
- error analysis, classification, [Error Analysis-Error Analysis](#)
- estimated versus actual probabilities, [The ROC Curve](#)
- estimators, [Clean the Data, Transformation Pipelines, Voting Classifiers](#)
- Euclidean norm, [Select a Performance Measure](#)
- event files, TensorBoard, [Using TensorBoard for Visualization](#)
- Example protobuf, [TensorFlow Protobufs](#)

- exclusive or (XOR) problem, [The Perceptron](#)
- exemplars, affinity propagation, [Other Clustering Algorithms](#)
- expectation-maximization (EM), [Gaussian Mixtures](#)
- experience replay, [The Difficulties of Training GANs](#)
- explainability, attention mechanisms, [Vision Transformers](#)
- explained variance ratio, [Explained Variance Ratio](#)
- explained variance, plotting, [Choosing the Right Number of Dimensions](#)-[Choosing the Right Number of Dimensions](#)
- exploding gradients, [The Vanishing/Exploding Gradients Problems](#)
 - (see also vanishing and exploding gradients)
- exploration policies, [Temporal Difference Learning](#), [Exploration Policies](#)
- exploration/exploitation dilemma, reinforcement learning, [Neural Network Policies](#)
- exponential linear unit (ELU), [ELU and SELU-GELU, Swish, and Mish](#)
- exponential scheduling, [Learning Rate Scheduling](#), [Learning Rate Scheduling](#)
- export_graphviz(), [Training and Visualizing a Decision Tree](#)
- extra-trees, random forest, [Extra-Trees](#)
- extremely randomized trees ensemble (extra-trees), [Extra-Trees](#)

F

- face-recognition classifier, [Multilabel Classification](#)
- false negatives, confusion matrix, [Confusion Matrices](#)
- false positive rate (FPR) or fall-out, [The ROC Curve](#)
- false positives, confusion matrix, [Confusion Matrices](#)
- fan-in/fan-out, [Glorot and He Initialization](#)
- fast-MCD, [Other Algorithms for Anomaly and Novelty Detection](#)
- FCNs (fully convolutional networks), [Fully Convolutional Networks](#)-[Fully Convolutional Networks](#), [Semantic Segmentation](#)
- feature engineering, [Irrelevant Features](#), [Clustering Algorithms: k-means and DBSCAN](#)
- feature extraction, [Unsupervised learning](#), [Irrelevant Features](#)
- feature maps, [Stacking Multiple Feature Maps](#)-[Stacking Multiple Feature Maps](#), [Implementing Convolutional Layers with Keras](#), [ResNet](#), [SENet](#)
- feature scaling, [Feature Scaling and Transformation](#)-[Feature Scaling and Transformation](#), [Gradient Descent](#), [Linear SVM Classification](#)

- feature selection, [Irrelevant Features](#), [Analyzing the Best Models and Their Errors](#), [Lasso Regression](#), [Feature Importance](#)
- feature vectors, [Select a Performance Measure](#), [Linear Regression](#), [Under the Hood of Linear SVM Classifiers](#)
- features, [Supervised learning](#)
- federated learning, [Running a Model in a Web Page](#)
- feedforward neural network (FNN), [The Multilayer Perceptron and Backpropagation](#), [Recurrent Neurons and Layers](#)
- fetch_openml(), [MNIST](#)
- fillna(), [Clean the Data](#)
- filters, convolutional layers, [Filters](#), [Implementing Convolutional Layers with Keras](#), [CNN Architectures](#), [Xception](#)
- first moment (mean of gradient), [Adam](#)
- first-order partial derivatives (Jacobians), [AdamW](#)
- fit()
 - and custom transformers, [Custom Transformers](#), [Transformation Pipelines](#)
 - data cleaning, [Clean the Data](#)
 - versus partial_fit(), [Stochastic Gradient Descent](#)
 - using only with training set, [Feature Scaling and Transformation](#)
- fitness function, [Model-based learning and a typical machine learning workflow](#)
- fit_transform(), [Clean the Data](#), [Feature Scaling and Transformation](#), [Custom Transformers](#), [Transformation Pipelines](#)
- fixed Q-value targets, [Fixed Q-value Targets](#)
- flat dataset, [Preparing the Data for Machine Learning Models](#)
- flowers dataset, [Pretrained Models for Transfer Learning](#)-[Pretrained Models for Transfer Learning](#)
- FNN (feedforward neural network), [The Multilayer Perceptron and Backpropagation](#), [Recurrent Neurons and Layers](#)
- folds, [Better Evaluation Using Cross-Validation](#), [MNIST](#), [Measuring Accuracy Using Cross-Validation](#), [Measuring Accuracy Using Cross-Validation](#)
- forecasting time series (see time series data)
- forget gate, LSTM, [LSTM cells](#)
- forward pass, in backpropagation, [The Multilayer Perceptron and Backpropagation](#)
- forward process, diffusion model, [Diffusion Models](#)-[Diffusion Models](#)
- FPR (false positive rate) or fall-out, [The ROC Curve](#)

- from_predictions(), [Error Analysis](#)
- full gradient descent, [Batch Gradient Descent](#)
- fully connected layer, [The Perceptron](#), [The Architecture of the Visual Cortex](#), [Memory Requirements](#)
- fully convolutional networks (FCNs), [Fully Convolutional Networks](#)-
[Fully Convolutional Networks](#), [Semantic Segmentation](#)
- function definition (FuncDef), [TF Functions and Concrete Functions](#)
- function graph (FuncGraph), [TF Functions and Concrete Functions](#)
- functional API, complex models with, [Building Complex Models Using the Functional API](#)-
[Building Complex Models Using the Functional API](#)
- FunctionTransformer, [Custom Transformers](#)
- F₁ score, [Precision and Recall](#)

G

- game play (see reinforcement learning)
- gamma (γ) value, [Gaussian RBF Kernel](#)
- GANs (see generative adversarial networks)
- gate controllers, LSTM, [LSTM cells](#)
- gated activation units, [WaveNet](#)
- gated recurrent unit (GRU) cell, [GRU cells](#)-[GRU cells](#), [Masking](#)
- Gaussian distribution, [Feature Scaling and Transformation](#), [Training and Cost Function](#), [Variational Autoencoders](#)-[Variational Autoencoders](#)
- Gaussian mixture model (GMM), [Gaussian Mixtures](#)-[Other Algorithms for Anomaly and Novelty Detection](#)
 - anomaly detection, [Using Gaussian Mixtures for Anomaly Detection](#)-[Using Gaussian Mixtures for Anomaly Detection](#)
 - Bayesian Gaussian mixtures, [Bayesian Gaussian Mixture Models](#)
 - fast-MCD, [Other Algorithms for Anomaly and Novelty Detection](#)
 - inverse_transform() with PCA, [Other Algorithms for Anomaly and Novelty Detection](#)
- isolation forest, [Other Algorithms for Anomaly and Novelty Detection](#)
- and k-means limitations, [Limits of k-means](#)
- local outlier factor, [Other Algorithms for Anomaly and Novelty Detection](#)
- one-class SVM, [Other Algorithms for Anomaly and Novelty Detection](#)
- selecting number of clusters, [Selecting the Number of Clusters](#)-
[Selecting the Number of Clusters](#)

- Gaussian process, [Fine-Tuning Neural Network Hyperparameters](#)
- Gaussian RBF kernel, [Custom Transformers](#), [Gaussian RBF Kernel-SVM Classes and Computational Complexity](#), [Kernelized SVMs](#)
- GBRT (gradient boosted regression trees), [Gradient Boosting](#), [Gradient Boosting-Histogram-Based Gradient Boosting](#)
- GCP (Google Cloud Platform), [Creating a Prediction Service on Vertex AI-Creating a Prediction Service on Vertex AI](#)
- GCS (Google Cloud Storage), [Creating a Prediction Service on Vertex AI](#)
- GD (see gradient descent)
- GELU, [GELU, Swish, and Mish-GELU, Swish, and Mish](#)
- generalization error, [Testing and Validating-Hyperparameter Tuning and Model Selection](#)
- generative adversarial networks (GANs), [Generative Adversarial Networks-StyleGANs](#)
 - deep convolutional, [Deep Convolutional GANs-Deep Convolutional GANs](#)
 - progressive growing of, [Progressive Growing of GANs-Progressive Growing of GANs](#)
 - StyleGANs, [StyleGANs-StyleGANs](#)
 - training difficulties, [The Difficulties of Training GANs-The Difficulties of Training GANs](#)
- generative autoencoders, [Variational Autoencoders](#)
- generative models, [Autoencoders, GANs, and Diffusion Models](#)
 - (see also Gaussian mixture model)
- generator, GAN, [Autoencoders, GANs, and Diffusion Models](#), [Generative Adversarial Networks-The Difficulties of Training GANs](#)
- genetic algorithm, [Policy Search](#)
- geodesic distance, [Other Dimensionality Reduction Techniques](#)
- geographic data, visualizing, [Visualizing Geographical Data-Visualizing Geographical Data](#)
- get_dummies(), [Handling Text and Categorical Attributes](#)
- get_feature_names_out(), [Custom Transformers](#)
- get_params(), [Custom Transformers](#)
- GINI impurity, [Making Predictions](#), [Gini Impurity or Entropy?](#)
- global average pooling layer, [Implementing Pooling Layers with Keras](#)
- global versus local minimum, gradient descent, [Gradient Descent](#)
- Glorot initialization, [Glorot and He Initialization-Glorot and He Initialization](#)
- GMM (see Gaussian mixture model)

- Google Cloud Platform (GCP), [Creating a Prediction Service on Vertex AI](#)-
[Creating a Prediction Service on Vertex AI](#)
- Google Cloud Storage (GCS), [Creating a Prediction Service on Vertex AI](#)
- Google Colab, [Running the Code Examples Using Google Colab](#)-
[Running the Code Examples Using Google Colab](#)
- Google Vertex AI (see Vertex AI)
- GoogLeNet, [GoogLeNet](#)-[GoogLeNet](#), [Xception](#)
- GPU implementation, [Mini-Batch Gradient Descent](#), [Using GPUs to Speed Up Computations](#)-[Parallel Execution Across Multiple Devices](#),
[Training at Scale Using the Distribution Strategies API](#)
 - getting your own GPU, [Getting Your Own GPU](#)-[Getting Your Own GPU](#)
 - managing RAM, [Managing the GPU RAM](#)-[Managing the GPU RAM](#)
 - operations handling, [Parallel Execution Across Multiple Devices](#)
 - parallel execution across multiple devices, [Parallel Execution Across Multiple Devices](#)-[Parallel Execution Across Multiple Devices](#)
 - placing operations and variables on devices, [Placing Operations and Variables on Devices](#)-[Placing Operations and Variables on Devices](#)
- gradient ascent, [Policy Search](#)
- gradient boosted regression trees (GBRT), [Gradient Boosting](#), [Gradient Boosting](#)-[Histogram-Based Gradient Boosting](#)
- gradient boosting, [Gradient Boosting](#)-[Gradient Boosting](#)
- gradient clipping, [Gradient Clipping](#)
- gradient descent (GD), [Training Models](#), [Gradient Descent](#)-[Mini-Batch Gradient Descent](#)
 - algorithm comparisons, [Mini-Batch Gradient Descent](#)
 - batch gradient descent, [Batch Gradient Descent](#)-[Batch Gradient Descent](#), [Ridge Regression](#)
 - local versus global minimum, [Gradient Descent](#)
 - mini-batch gradient descent, [Mini-Batch Gradient Descent](#)-[Mini-Batch Gradient Descent](#)
 - minimizing hinge loss, [Under the Hood of Linear SVM Classifiers](#)
 - versus momentum optimization, [Momentum](#)
 - with optimizers, [Faster Optimizers](#)-[AdamW](#)
 - shuffling data, [Shuffling the Data](#)
 - stochastic gradient descent, [Stochastic Gradient Descent](#)-[Stochastic Gradient Descent](#)
- gradient tree boosting, [Gradient Boosting](#)

- gradients
 - autodiff for computing, [Computing Gradients Using Autodiff](#)-
[Computing Gradients Using Autodiff](#)
 - bandwidth saturation issue, [Bandwidth saturation](#)
 - PG algorithm, [Policy Search](#), [Policy Gradients](#)-[Policy Gradients](#)
 - stale, [Asynchronous updates](#)
 - unstable (see vanishing and exploding gradients)
- graph mode, [AutoGraph and Tracing](#)
- graphical processing units (see GPU implementation)
- graphs and functions, TensorFlow, [A Quick Tour of TensorFlow](#),
[TensorFlow Functions and Graphs](#)-[TF Function Rules](#), [TensorFlow Graphs](#)-[Using TF Functions with Keras \(or Not\)](#)
- Graphviz, [Training and Visualizing a Decision Tree](#)
- greedy algorithm, CART as, [The CART Training Algorithm](#)
- greedy decoding, [Generating Fake Shakespearean Text](#)
- greedy layer-wise pretraining, [Unsupervised Pretraining](#), [Training One Autoencoder at a Time](#), [Progressive Growing of GANs](#)
- grid search, [Grid Search](#)-[Grid Search](#)
- GridSearchCV, [Grid Search](#)-[Grid Search](#)
- gRPC API, querying through, [Querying TF Serving through the gRPC API](#)
- GRU (gated recurrent unit) cell, [GRU cells](#)-[GRU cells](#), [Masking](#)

H

- hard clustering, [k-means](#)
- hard margin classification, [Soft Margin Classification](#), [Under the Hood of Linear SVM Classifiers](#)
- hard voting classifiers, [Voting Classifiers](#)
- harmonic mean, [Precision and Recall](#)
- hashing collision, [The StringLookup Layer](#)
- Hashing layer, [The Hashing Layer](#)
- hashing trick, [The StringLookup Layer](#)
- HDBSCAN (hierarchical DBSCAN), [DBSCAN](#)
- He initialization, [Glorot and He Initialization](#)-[Glorot and He Initialization](#)
- Heaviside step function, [The Perceptron](#)
- heavy tail, feature distribution, [Feature Scaling and Transformation](#)
- Hebb's rule, [The Perceptron](#)
- Hebbian learning, [The Perceptron](#)

- Hessians, [AdamW](#)
- hidden layers
 - neurons per layer, [Number of Neurons per Hidden Layer](#)
 - number of, [Number of Hidden Layers](#)
 - stacked autoencoders, [Stacked Autoencoders-Training One Autoencoder at a Time](#)
- hierarchical clustering, [Unsupervised learning](#)
- hierarchical DBSCAN (HDBSCAN), [DBSCAN](#)
- high variance, with decision trees, [Decision Trees Have a High Variance](#)
- hinge loss function, [Under the Hood of Linear SVM Classifiers](#)
- histogram-based gradient boosting (HGB), [Histogram-Based Gradient Boosting-Histogram-Based Gradient Boosting](#)
- histograms, [Take a Quick Look at the Data Structure](#)
- hold-out sets, [Testing and Validating](#)
- holdout validation, [Hyperparameter Tuning and Model Selection](#)
- housing dataset, [Working with Real Data-Check the Assumptions](#)
- Huber loss, [Regression MLPs](#), [Custom Loss Functions](#), [Custom Metrics](#), [Forecasting Using a Linear Model](#)
- Hugging Face, [Hugging Face's Transformers Library-Hugging Face's Transformers Library](#)
- Hungarian algorithm, [Object Tracking](#)
- Hyperband tuner, [Fine-Tuning Neural Network Hyperparameters](#)
- hyperbolic tangent (htan), [The Multilayer Perceptron and Backpropagation](#), [Fighting the Unstable Gradients Problem](#)
- hyperparameters, [Overfitting the Training Data](#), [Fine-Tuning Neural Network Hyperparameters-Learning Rate, Batch Size, and Other Hyperparameters](#)
 - activation function, [Learning Rate, Batch Size, and Other Hyperparameters](#)
 - batch size, [Learning Rate, Batch Size, and Other Hyperparameters](#)
 - CART algorithm, [The CART Training Algorithm](#)
 - convolutional layers, [Implementing Convolutional Layers with Keras](#)
 - in custom transformations, [Custom Transformers](#)
 - decision tree, [Gradient Boosting](#)
 - dimensionality reduction, [Choosing the Right Number of Dimensions](#)
 - gamma (γ) value, [Gaussian RBF Kernel](#)

- GAN challenges, [The Difficulties of Training GANs](#)
- Keras Tuner, [Hyperparameter Tuning on Vertex AI](#)
- learning rate, [Gradient Descent](#), [Learning Rate, Batch Size, and Other Hyperparameters](#)
- momentum β , [Momentum](#)
- Monte Carlo samples, [Monte Carlo \(MC\) Dropout](#)
- neurons per hidden layer, [Number of Neurons per Hidden Layer](#)
- and normalization, [Feature Scaling and Transformation](#)
- number of hidden layers, [Number of Hidden Layers](#)
- number of iterations, [Learning Rate, Batch Size, and Other Hyperparameters](#)
- optimizer, [Learning Rate, Batch Size, and Other Hyperparameters](#)
- PG algorithms, [Policy Gradients](#)
- preprocessor and model interaction, [Grid Search](#)
- randomized search, [Fine-Tuning Neural Network Hyperparameters](#)
- saving along with model, [Custom Activation Functions, Initializers, Regularizers, and Constraints](#)
- SGDClassifier, [SVM Classes and Computational Complexity](#)
- subsample, [Gradient Boosting](#)
- SVM classifiers with polynomial kernel, [Polynomial Kernel](#)
- tolerance (ε), [SVM Classes and Computational Complexity](#)
- tuning of, [Hyperparameter Tuning and Model Selection](#)-
[Hyperparameter Tuning and Model Selection](#), [Grid Search-Grid Search](#), [Evaluate Your System on the Test Set](#), [Training and evaluating the model](#), [Hyperparameter Tuning on Vertex AI](#)-
[Hyperparameter Tuning on Vertex AI](#)
- hypothesis, [Select a Performance Measure](#)
- hypothesis boosting (see boosting)
- hypothesis function, [Linear Regression](#)

I

- identity matrix, [Ridge Regression](#)
- IID (see independent and identically distributed)
- image generation, [Semantic Segmentation](#), [StyleGANs](#), [Diffusion Models](#)
- image segmentation, [Clustering Algorithms: k-means and DBSCAN](#),
[Using Clustering for Image Segmentation](#)-[Using Clustering for Image Segmentation](#)

- images, classifying and generating, [Examples of Applications](#)
 - autoencoders (see autoencoders)
 - CNNs (see convolutional neural networks)
 - diffusion models, [Diffusion Models-Diffusion Models](#)
 - generating with GANs, [Generative Adversarial Networks-StyleGANs](#)
 - implementing MLPs, [Building an Image Classifier Using the Sequential API-Using the model to make predictions](#)
 - labels, [Classification and Localization](#)
 - loading and preprocessing data, [Image Preprocessing Layers](#)
 - representative images, [Using Clustering for Semi-Supervised Learning](#)
 - semantic segmentation, [Using Clustering for Image Segmentation](#)
 - tuning hyperparameters, [Fine-Tuning Neural Network Hyperparameters](#)
- importance sampling (IS), [Prioritized Experience Replay](#)
- impurity measures, [Making Predictions](#), [Gini Impurity or Entropy?](#)
- imputation, [Clean the Data](#)
- incremental learning, [Online learning](#)
- incremental PCA (IPCA), [Incremental PCA-Incremental PCA](#)
- independent and identically distributed (IID), training instances as, [Stochastic Gradient Descent](#)
- inductive bias, [Vision Transformers](#)
- inertia, model, [Centroid initialization methods](#), [Accelerated k-means and mini-batch k-means](#)
- inference, [Model-based learning and a typical machine learning workflow](#)
- info(), [Take a Quick Look at the Data Structure](#)
- information theory, [Softmax Regression](#), [Gini Impurity or Entropy?](#)
- inliers, [Unsupervised Learning Techniques](#)
- input and output sequences, RNNs, [Input and Output Sequences-Input and Output Sequences](#)
- input gate, LSTM, [LSTM cells](#)
- input layer, neural network, [The Perceptron](#), [Creating the model using the sequential API](#), [Building Complex Models Using the Functional API](#)
 - (see also hidden layers)
- input signature, [TF Functions and Concrete Functions](#)
- instance segmentation, [Using Clustering for Image Segmentation](#), [Semantic Segmentation](#)

- instance-based learning, [Instance-based learning](#), [Model-based learning and a typical machine learning workflow](#)
- inter-op thread pool, [Parallel Execution Across Multiple Devices](#)
- intercept term constant, [Linear Regression](#)
- interleaving lines from multiple files, [Interleaving Lines from Multiple Files](#)-[Interleaving Lines from Multiple Files](#)
- interpretable ML, [Making Predictions](#)
- invariance, max pooling layer, [Pooling Layers](#)
- inverse_transform(), [Feature Scaling and Transformation](#), [Custom Transformers](#), [PCA for Compression](#), [Other Algorithms for Anomaly and Novelty Detection](#)
- IPCA (incremental PCA), [Incremental PCA](#)-[Incremental PCA](#)
- iris dataset, [Decision Boundaries](#)
- irreducible error, [Learning Curves](#)
- IS (importance sampling), [Prioritized Experience Replay](#)
- isolation forest, [Other Algorithms for Anomaly and Novelty Detection](#)
- Isomap, [Other Dimensionality Reduction Techniques](#)
- isotropic noise, [Diffusion Models](#)
- IterativeImputer, [Clean the Data](#)

J

- Jacobians, [AdamW](#)
- joblib library, [Launch, Monitor, and Maintain Your System](#)-[Launch, Monitor, and Maintain Your System](#)
- JSON Lines, [Running Batch Prediction Jobs on Vertex AI](#), [Running Batch Prediction Jobs on Vertex AI](#)
- Jupyter, [Running the Code Examples Using Google Colab](#)

K

- k-fold cross-validation, [Better Evaluation Using Cross-Validation](#), [Measuring Accuracy Using Cross-Validation](#), [Measuring Accuracy Using Cross-Validation](#)
- k-means algorithm, [Custom Transformers](#), [k-means-Limits of k-means](#)
 - accelerated k-means, [Accelerated k-means and mini-batch k-means](#)
 - centroid initialization methods, [Centroid initialization methods](#)-[Centroid initialization methods](#)
 - finding optimal number of clusters, [Finding the optimal number of clusters](#)-[Finding the optimal number of clusters](#)

- limitations of, [Limits of k-means](#)
- mini-batch k-means, [Accelerated k-means and mini-batch k-means](#)
- workings of, [The k-means algorithm](#)-[The k-means algorithm](#)
- k-means++, [Centroid initialization methods](#)
- k-nearest neighbors regression, [Model-based learning and a typical machine learning workflow](#)
- Kaiming initialization, [Glorot and He Initialization](#)
- Kalman Filters, [Object Tracking](#)
- Keras API, [Objective and Approach](#), [A Quick Tour of TensorFlow](#)
 - (see also artificial neural networks)
 - and accessing TensorFlow API directly, [Image Preprocessing Layers](#)
 - activation function support, [Leaky ReLU](#), [ELU](#) and [SELU](#), [GELU](#), [Swish](#), and [Mish](#)
 - convolutional layer implementation, [Implementing Convolutional Layers with Keras](#)-[Implementing Convolutional Layers with Keras](#)
 - custom functions in, [TensorFlow Functions and Graphs](#)
 - gradient clipping, [Gradient Clipping](#)
 - image preprocessing layers, [Image Preprocessing Layers](#)
 - implementing MLPs with, [Implementing MLPs with Keras](#)-[Using TensorBoard for Visualization](#)
 - initialization handling, [Glorot and He Initialization](#)
 - initializers, [Creating the model using the sequential API](#)
 - layer preprocessing, [Keras Preprocessing Layers](#)-[Using Pretrained Language Model Components](#)
 - learning rate scheduling, [Learning Rate Scheduling](#)-[Learning Rate Scheduling](#)
 - loading a dataset, [Building an Image Classifier Using the Sequential API](#)
 - PG algorithm, [Policy Gradients](#)-[Policy Gradients](#)
 - pool layer implementation, [Implementing Pooling Layers with Keras](#)-[Implementing Pooling Layers with Keras](#)
 - pretrained CNN models, [Using Pretrained Models from Keras](#)-[Using Pretrained Models from Keras](#)
 - ResNet-34 CNN with, [Implementing a ResNet-34 CNN Using Keras](#)
 - saving models, [Exporting SavedModels](#), [Deploying a Model to a Mobile or Embedded Device](#)
 - stacked encoder implementation, [Implementing a Stacked Autoencoder Using Keras](#)

- tf.data API dataset, [Using the Dataset with Keras](#)-[Using the Dataset with Keras](#)
- tf.keras library, [Using Keras to load the dataset](#)
- tf.keras.activations.get(), [Custom Layers](#)
- tf.keras.activations.relu(), [Creating the model using the sequential API](#)
- tf.keras.applications module, [Using Pretrained Models from Keras](#)
- tf.keras.applications.xception.preprocess_input(), [Pretrained Models for Transfer Learning](#)
- tf.keras.backend module, [Tensors and Operations](#)
- tf.keras.callbacks.EarlyStopping, [Using Callbacks](#)
- tf.keras.callbacks.LearningRateScheduler, [Learning Rate Scheduling](#)
- tf.keras.callbacks.ModelCheckpoint, [Using Callbacks](#)
- tf.keras.callbacks.TensorBoard, [Using TensorBoard for Visualization](#), [Masking](#)
- tf.keras.datasets.imdb.load_data(), [Sentiment Analysis](#)
- tf.keras.initializers.VarianceScaling, [Glorot and He Initialization](#)
- tf.keras.layers.ActivityRegularization, [Sparse Autoencoders](#)
- tf.keras.layers.AdditiveAttention, [Attention Mechanisms](#)
- tf.keras.layers.Attention, [Attention Mechanisms](#)
- tf.keras.layers.AvgPool2D, [Implementing Pooling Layers with Keras](#)
- tf.keras.layers.BatchNormalization, [Batch Normalization](#),
[Implementing batch normalization with Keras](#)-[Implementing batch normalization with Keras](#), [Fighting the Unstable Gradients Problem](#)
- tf.keras.layers.Bidirectional, [Bidirectional RNNs](#)
- tf.keras.layers.CategoryEncoding, [The CategoryEncoding Layer](#)
- tf.keras.layers.CenterCrop, [Image Preprocessing Layers](#)
- tf.keras.layers.Concatenate, [Building Complex Models Using the Functional API](#), [Building Complex Models Using the Functional API](#), [GoogLeNet](#)
- tf.keras.layers.Conv1D, [Semantic Segmentation](#), [Masking](#)
- tf.keras.layers.Conv2D, [Implementing Convolutional Layers with Keras](#)
- tf.keras.layers.Conv2DTranspose, [Semantic Segmentation](#)
- tf.keras.layers.Conv3D, [Semantic Segmentation](#)
- tf.keras.layers.Dense, [Building Complex Models Using the Functional API](#), [Building Complex Models Using the Functional API](#), [GoogLeNet](#)

[Custom Layers](#)-[Custom Layers](#), [Encoding Categorical Features Using Embeddings](#), [Tying Weights](#), [Variational Autoencoders](#)

- tf.keras.layers.Discretization, [The Discretization Layer](#)
- tf.keras.layers.Dropout, [Dropout](#)
- tf.keras.layers.Embedding, [Encoding Categorical Features Using Embeddings](#), [Building and Training the Char-RNN Model](#), [Sentiment Analysis](#), [Positional encodings](#)
- tf.keras.layers.GlobalAvgPool2D, [Implementing Pooling Layers with Keras](#)
- tf.keras.layers.GRU, [GRU cells](#)
- tf.keras.layers.GRUCell, [GRU cells](#)
- tf.keras.layers.Hashing, [The Hashing Layer](#)
- tf.keras.layers.Input, [Building Complex Models Using the Functional API](#)
- tf.keras.layers.Lambda, [Custom Layers](#), [Building and Training the Char-RNN Model](#)
- tf.keras.layers.LayerNormalization, [Fighting the Unstable Gradients Problem](#)
- tf.keras.layers.LeakyReLU, [Leaky ReLU](#)
- tf.keras.layers.LSTM, [LSTM cells](#)
- tf.keras.layers.Masking, [Masking](#)
- tf.keras.layers.MaxPool2D, [Implementing Pooling Layers with Keras](#)
- tf.keras.layers.MultiHeadAttention, [Multi-head attention](#)
- tf.keras.layers.Normalization, [Building a Regression MLP Using the Sequential API](#), [Building Complex Models Using the Functional API](#), [The Normalization Layer](#)-[The Normalization Layer](#)
- tf.keras.layers.PReLU, [Leaky ReLU](#)
- tf.keras.layers.Rescaling, [Image Preprocessing Layers](#)
- tf.keras.layers.Resizing, [Image Preprocessing Layers](#), [Using Pretrained Models from Keras](#)
- tf.keras.layers.RNN, [LSTM cells](#)
- tf.keras.layers.SeparableConv2D, [Xception](#)
- tf.keras.layers.StringLookup, [The StringLookup Layer](#)
- tf.keras.layers.TextVectorization, [Text Preprocessing](#)-[Text Preprocessing](#), [Creating the Training Dataset](#), [Building and Training the Char-RNN Model](#), [Sentiment Analysis](#), [Sentiment Analysis](#)-[An Encoder-Decoder Network for Neural Machine Translation](#)

- tf.keras.layers.TimeDistributed, [Forecasting Using a Sequence-to-Sequence Model](#)
- tf.keras.losses.Huber, [Custom Loss Functions](#)
- tf.keras.losses.kullback_leibler_divergence(), [Sparse Autoencoders](#)
- tf.keras.losses.Loss, [Saving and Loading Models That Contain Custom Components](#)
- tf.keras.losses.sparse_categorical_crossentropy(), [Compiling the model](#), [CNN Architectures](#), [Building and Training the Char-RNN Model](#), [An Encoder–Decoder Network for Neural Machine Translation](#), [Hugging Face’s Transformers Library](#)
- tf.keras.metrics.MeanIoU, [Classification and Localization](#)
- tf.keras.metrics.Metric, [Custom Metrics](#)
- tf.keras.metrics.Precision, [Custom Metrics](#)
- tf.keras.Model, [Building Complex Models Using the Functional API](#)
- tf.keras.models.clone_model(), [Using the Subclassing API to Build Dynamic Models](#), [Transfer Learning with Keras](#)
- tf.keras.models.load_model(), [Saving and Restoring a Model](#), [Saving and Loading Models That Contain Custom Components-Saving and Loading Models That Contain Custom Components](#), [Custom Models](#), [Training at Scale Using the Distribution Strategies API](#)
- tf.keras.optimizers.Adam, [Building a Regression MLP Using the Sequential API](#), [AdamW](#)
- tf.keras.optimizers.Adamax, [AdamW](#)
- tf.keras.optimizers.experimental.AdamW, [AdamW](#)
- tf.keras.optimizers.Nadam, [AdamW](#)
- tf.keras.optimizers.schedules, [Learning Rate Scheduling](#)
- tf.keras.optimizers.SGD, [Momentum](#)
- tf.keras.preprocessing.image.ImageDataGenerator, [Pretrained Models for Transfer Learning](#)
- tf.keras.regularizers.l1_l20, [ℓ₁ and ℓ₂ Regularization](#)
- tf.keras.Sequential, [Creating the model using the sequential API](#), [Building a Regression MLP Using the Sequential API](#), [CNN Architectures](#)
- tf.keras.utils.get_file(), [Creating the Training Dataset](#)
- tf.keras.utils.set_random_seed(), [Creating the model using the sequential API](#)
- tf.keras.utils.timeseries_dataset_from_array(), [Preparing the Data for Machine Learning Models](#), [Preparing the Data for Machine Learning Models](#)

- `tf.keras.utils.to_categorical()`, [Compiling the model](#)
- time series forecasting for RNN, [Forecasting Using a Simple RNN](#)-
[Forecasting Using a Deep RNN](#)
- transfer learning with, [Transfer Learning with Keras](#)-[Transfer Learning with Keras](#)
- using TF functions (or not), [Using TF Functions with Keras \(or Not\)](#)
- Keras session, [Creating the model using the sequential API](#)
- Keras Tuner, [Hyperparameter Tuning on Vertex AI](#)
- kernel trick, [Polynomial Kernel-SVM Classes and Computational Complexity](#), [Kernelized SVMs](#)-[Kernelized SVMs](#)
- kernelized SVMs, [Kernelized SVMs](#)-[Kernelized SVMs](#)
- kernels (convolution kernels), [Filters](#), [CNN Architectures](#)
- kernels (runtimes), [Running the Code Examples Using Google Colab](#)
- KL (Kullback-Leibler) divergence, [Softmax Regression](#), [Sparse Autoencoders](#)
- `KLDivergenceRegularizer`, [Sparse Autoencoders](#)
- KMeans, [Custom Transformers](#)
- KNeighborsClassifier, [Multilabel Classification](#), [Multioutput Classification](#)
- KNNImputer, [Clean the Data](#)
- Kullback-Leibler (KL) divergence, [Softmax Regression](#), [Sparse Autoencoders](#)

L

- label propagation, [Using Clustering for Semi-Supervised Learning](#)-
[Using Clustering for Semi-Supervised Learning](#)
- labels, [Frame the Problem](#)
 - in clustering, [k-means](#)
 - image classification, [Classification and Localization](#)
 - supervised learning, [Supervised learning](#)
 - unlabeled data issue, [Unsupervised Pretraining Using Stacked Autoencoders](#)
- landmarks, [Similarity Features](#)
- language models, [Generating Shakespearean Text Using a Character RNN](#)
 - (see also natural language processing)
- large margin classification, [Linear SVM Classification](#)
- Lasso, [Lasso Regression](#)
- lasso regression, [Lasso Regression](#)-[Lasso Regression](#)

- latent diffusion models, [Diffusion Models](#)
- latent loss, [Variational Autoencoders](#)
- latent representation of inputs, [Vision Transformers](#), [Autoencoders](#), [GANs, and Diffusion Models](#), [Efficient Data Representations](#), [Deep Convolutional GANs](#)
- latent space, [Variational Autoencoders](#)
- law of large numbers, [Voting Classifiers](#)
- layer normalization, [Exercises](#), [Processing Sequences Using RNNs and CNNs](#), [Fighting the Unstable Gradients Problem](#)
- LDA (linear discriminant analysis), [Other Dimensionality Reduction Techniques](#)
- leaf node, decision tree, [Making Predictions](#), [Estimating Class Probabilities](#), [Regularization Hyperparameters](#)
- leakyReLU, [Leaky ReLU-Leaky ReLU](#)
- learning curves, overfit or underfit analysis, [Learning Curves](#)-[Learning Curves](#)
- learning rate, [Online learning](#), [Gradient Descent](#), [Batch Gradient Descent](#), [Stochastic Gradient Descent](#), [Learning Rate](#), [Batch Size, and Other Hyperparameters](#), [Q-Learning](#)
- learning rate schedules, [Learning Rate Scheduling](#)-[Learning Rate Scheduling](#)
- learning schedules, [Stochastic Gradient Descent](#)
- LeCun initialization, [Glorot and He Initialization](#)
- LeNet-5, [The Architecture of the Visual Cortex](#), [LeNet-5](#)
- Levenshtein distance, [Gaussian RBF Kernel](#)
- liblinear library, [SVM Classes and Computational Complexity](#)
- libsvm library, [SVM Classes and Computational Complexity](#)
- life satisfaction dataset, [Model-based learning and a typical machine learning workflow](#)
- likelihood function, [Selecting the Number of Clusters](#)-[Selecting the Number of Clusters](#)
- linear discriminant analysis (LDA), [Other Dimensionality Reduction Techniques](#)
- linear models
 - forecasting time series, [Forecasting Using a Linear Model](#)
 - linear regression (see linear regression)
 - regularized, [Regularized Linear Models](#)-[Early Stopping](#)
 - SVM, [Linear SVM Classification](#)-[Soft Margin Classification](#)

- training and running example, [Model-based learning and a typical machine learning workflow](#)
- linear regression, [Model-based learning and a typical machine learning workflow](#)-[Model-based learning and a typical machine learning workflow](#), [Linear Regression](#)-[Mini-Batch Gradient Descent](#)
 - comparison of algorithms, [Mini-Batch Gradient Descent](#)
 - computational complexity, [Computational Complexity](#)
 - gradient descent in, [Gradient Descent](#)-[Mini-Batch Gradient Descent](#)
 - learning curves in, [Learning Curves](#)-[Learning Curves](#)
 - Normal equation, [The Normal Equation](#)-[The Normal Equation](#)
 - regularizing models (see regularization)
 - ridge regression, [Ridge Regression](#)-[Ridge Regression](#), [Elastic Net Regression](#)
 - training set evaluation, [Train and Evaluate on the Training Set](#)
 - using stochastic gradient descent, [Stochastic Gradient Descent](#)
- linear SVM classification, [Linear SVM Classification](#)-[Soft Margin Classification](#), [Under the Hood of Linear SVM Classifiers](#)-[Kernelized SVMs](#)
- linear threshold units (LTUs), [The Perceptron](#)
- linearly separable, SVM classes, [Linear SVM Classification](#)-[Linear SVM Classification](#), [Similarity Features](#)
- LinearRegression, [Clean the Data](#), [Feature Scaling and Transformation](#), [Train and Evaluate on the Training Set](#), [The Normal Equation](#), [Computational Complexity](#), [Polynomial Regression](#)
- LinearSVC, [Soft Margin Classification](#), [SVM Classes and Computational Complexity](#), [Under the Hood of Linear SVM Classifiers](#)
- LinearSVR, [SVM Regression](#)-[SVM Regression](#)
- Lipschitz continuous, derivative as, [Gradient Descent](#)
- LLE (locally linear embedding), [LLE](#)-[LLE](#)
- loading and preprocessing data, [Loading and Preprocessing Data with TensorFlow](#)-[The TensorFlow Datasets Project](#)
 - image preprocessing layers, [Image Preprocessing Layers](#)
 - layer preprocessing in Keras, [Loading and Preprocessing Data with TensorFlow](#), [Keras Preprocessing Layers](#)-[Using Pretrained Language Model Components](#)
 - tf.data API, [Loading and Preprocessing Data with TensorFlow](#)-[Using the Dataset with Keras](#)
 - TFDS project, [The TensorFlow Datasets Project](#)-[The TensorFlow Datasets Project](#)

- TFRecord format, [Loading and Preprocessing Data with TensorFlow](#), [The TFRecord Format-Handling Lists of Lists Using the SequenceExample Protobuf](#)
- local outlier factor (LOF), [Other Algorithms for Anomaly and Novelty Detection](#)
- local receptive field, [The Architecture of the Visual Cortex](#)
- local response normalization (LRN), [AlexNet](#)
- local versus global minimum, gradient descent, [Gradient Descent](#)
- locality sensitive hashing (LSH), [Random Projection](#)
- localization, CNNs, [Classification and Localization-Object Detection](#)
- locally linear embedding (LLE), [LLE-LLE](#)
- LOF (local outlier factor), [Other Algorithms for Anomaly and Novelty Detection](#)
- log loss, [Training and Cost Function](#)
- log-odds function, [Estimating Probabilities](#)
- log-transformer, [Custom Transformers](#)
- logical GPU device, [Managing the GPU RAM](#)
- logistic function, [Estimating Probabilities](#)
- logistic regression, [Supervised learning](#), [Feature Scaling and Transformation](#), [Logistic Regression-Softmax Regression](#)
 - decision boundaries illustration, [Decision Boundaries-Decision Boundaries](#)
 - estimating probabilities, [Estimating Probabilities-Estimating Probabilities](#)
 - softmax regression model, [Softmax Regression-Softmax Regression](#)
 - training and cost function, [Training and Cost Function-Training and Cost Function](#)
- LogisticRegression, [Decision Boundaries](#), [Softmax Regression](#)
- logit function, [Estimating Probabilities](#)
- long sequences, training RNN on, [Handling Long Sequences-WaveNet](#)
 - short-term memory problem, [Tackling the Short-Term Memory Problem-WaveNet](#)
 - unstable gradients problem, [Fighting the Unstable Gradients Problem-Fighting the Unstable Gradients Problem](#)
- long short-term memory (LSTM) cell, [LSTM cells-LSTM cells](#), [Masking](#), [An Encoder-Decoder Network for Neural Machine Translation](#), [Bidirectional RNNs](#)
- loss functions

- based on model internals, [Losses and Metrics Based on Model Internals](#)-[Losses and Metrics Based on Model Internals](#)
- custom, [Custom Loss Functions](#)
- versus metrics, [Custom Metrics](#)
- output, [Building Complex Models Using the Functional API](#)
- LRN (local response normalization), [AlexNet](#)
- LSH (locality sensitive hashing), [Random Projection](#)
- LSTM (long short-term memory) cell, [LSTM cells](#)-[LSTM cells](#), [Masking](#),
[An Encoder–Decoder Network for Neural Machine Translation](#),
[Bidirectional RNNs](#)
- LTUs (linear threshold units), [The Perceptron](#)
- Luong attention, [Attention Mechanisms](#)

M

- machine learning (ML), [What Is Machine Learning?](#)
 - application/technique examples, [Examples of Applications](#)-[Examples of Applications](#)
 - challenges of, [Main Challenges of Machine Learning](#)-[Stepping Back](#)
 - notations, [Select a Performance Measure](#)-[Select a Performance Measure](#)
 - project checklist, [Look at the Big Picture](#)
 - (see also end-to-end ML project exercise)
 - reasons for using, [Why Use Machine Learning?](#)-[Why Use Machine Learning?](#)
 - resources on, [Other Resources](#)-[Other Resources](#)
 - spam filter example, [The Machine Learning Landscape](#)-[Why Use Machine Learning?](#)
 - testing and validating, [Testing and Validating](#)-[Data Mismatch](#)
 - types of systems, [Types of Machine Learning Systems](#)-[Model-based learning and a typical machine learning workflow](#)
- MAE (mean absolute error), [Select a Performance Measure](#)
- majority-vote predictions, [Exercises](#)
- make_column_selector(), [Transformation Pipelines](#)
- make_column_transformer(), [Transformation Pipelines](#)
- make_pipeline(), [Transformation Pipelines](#), [Learning Curves](#)
- Manhattan norm, [Select a Performance Measure](#)
- manifold hypothesis, [Manifold Learning](#)
- manifold learning, dimension reduction, [Manifold Learning](#)-[Manifold Learning](#)

- MAP (maximum a-posteriori) estimation, [Selecting the Number of Clusters](#)
- mAP (mean average precision), [You Only Look Once](#)
- MAPE (mean absolute percentage error), [Forecasting a Time Series](#)
- mapping network, StyleGANs, [StyleGANs](#)
- MapReduce, [Frame the Problem](#)
- margin violations, [Soft Margin Classification](#), [Under the Hood of Linear SVM Classifiers](#)
- Markov chains, [Markov Decision Processes](#)
- Markov decision processes (MDPs), [Markov Decision Processes](#)-[Markov Decision Processes](#), [Exploration Policies](#)
- mask R-CNN, [Semantic Segmentation](#)
- mask tensor, [Masking](#)
- masked language model (MLM), [An Avalanche of Transformer Models](#)
- masked multi-head attention layer, [Attention Is All You Need: The Original Transformer Architecture](#)
- masking, [Masking-Masking](#), [Multi-head attention](#)
- Matching Engine service, Vertex AI, [Creating a Prediction Service on Vertex AI](#)
- Matplotlib, [Running the Code Examples Using Google Colab](#)
- max pooling layer, [Pooling Layers](#), [CNN Architectures](#)
- max-norm regularization, [Max-Norm Regularization](#)
- maximization step, Gaussian mixtures, [Gaussian Mixtures](#)
- maximum a-posteriori (MAP) estimation, [Selecting the Number of Clusters](#)
- maximum likelihood estimate (MLE), [Selecting the Number of Clusters](#)
- MC (Monte Carlo) dropout regularization, [Monte Carlo \(MC\) Dropout](#)-[Monte Carlo \(MC\) Dropout](#)
- MCTS (Monte Carlo tree search), [Overview of Some Popular RL Algorithms](#)
- MDPs (Markov decision processes), [Markov Decision Processes](#)-[Markov Decision Processes](#), [Exploration Policies](#)
- MDS (multidimensional scaling), [Other Dimensionality Reduction Techniques](#)
- mean absolute error (MAE), [Select a Performance Measure](#)
- mean absolute percentage error (MAPE), [Forecasting a Time Series](#)
- mean average precision (mAP), [You Only Look Once](#)
- mean squared error (MSE), [Linear Regression](#), [Variational Autoencoders](#)

- mean-shift, clustering algorithms, [Other Clustering Algorithms](#)
- mean_squared_error(), [Train and Evaluate on the Training Set](#)
- measure of similarity, [Instance-based learning](#)
- memory bandwidth, GPU card, [Prefetching](#)
- memory cells (cells), RNNs, [Memory Cells](#), [Tackling the Short-Term Memory Problem-WaveNet](#)
- memory requirements, convolutional layers, [Memory Requirements](#)
- Mercer's theorem, [Kernelized SVMs](#)
- meta learner, [Stacking](#)
- metagraphs, SavedModel, [Exporting SavedModels](#)
- min-max scaling, [Feature Scaling and Transformation](#)
- mini-batch discrimination, [The Difficulties of Training GANs](#)
- mini-batch gradient descent, [Mini-Batch Gradient Descent-Mini-Batch Gradient Descent](#), [Early Stopping](#)
- mini-batch k-means, [Accelerated k-means and mini-batch k-means](#)
- mini-batches, [Online learning](#), [Progressive Growing of GANs](#)
- MinMaxScaler, [Feature Scaling and Transformation](#)
- mirrored strategy, data parallelism, [Data parallelism using the mirrored strategy](#), [Training at Scale Using the Distribution Strategies API](#), [Training a Model on a TensorFlow Cluster](#)
- Mish activation function, [GELU, Swish, and Mish](#)
- mixing regularization, StyleGAN, [StyleGANs](#)
- ML (see machine learning)
- ML Operations (MLOps), [Launch, Monitor, and Maintain Your System](#)
- MLE (maximum likelihood estimate), [Selecting the Number of Clusters](#)
- MLM (masked language model), [An Avalanche of Transformer Models](#)
- MLPs (see multilayer perceptrons)
- MNIST dataset, [MNIST-MNIST](#)
- mobile device, deploying model to, [Deploying a Model to a Mobile or Embedded Device](#)-[Deploying a Model to a Mobile or Embedded Device](#)
- mode collapse, [Vision Transformers](#), [The Difficulties of Training GANs](#)
- model parallelism, [Model Parallelism-Model Parallelism](#)
- model parameters, [Model-based learning and a typical machine learning workflow](#)
 - early stopping regularization, [Training and Cost Function](#)
 - in gradient descent, [Gradient Descent](#), [Batch Gradient Descent](#)
 - linear SVM classifier mechanics, [Under the Hood of Linear SVM Classifiers](#)
 - and variable updating, [Variables](#)

- weight matrix shape, [Creating the model using the sequential API](#)
- model rot, [Batch learning](#)
- model selection, [Model-based learning and a typical machine learning workflow](#), [Hyperparameter Tuning and Model Selection-Hyperparameter Tuning and Model Selection](#)
- model server (see TensorFlow Serving)
- model warmup, [Deploying a new model version](#)
- model-based learning, [Model-based learning and a typical machine learning workflow](#)-[Model-based learning and a typical machine learning workflow](#)
- modes, [Feature Scaling and Transformation](#)
- momentum optimization, [Momentum](#)
- momentum β , [Momentum](#)
- Monte Carlo (MC) dropout, [Monte Carlo \(MC\) Dropout](#)-[Monte Carlo \(MC\) Dropout](#)
- Monte Carlo tree search (MCTS), [Overview of Some Popular RL Algorithms](#)
- MSE (mean squared error), [Linear Regression](#), [Variational Autoencoders](#)
- multi-head attention layer, [Attention Is All You Need: The Original Transformer Architecture](#), [Multi-head attention](#)-[Multi-head attention](#)
- multi-hot encoding, [The CategoryEncoding Layer](#)
- multiclass (multinomial) classification, [Multiclass Classification](#)-[Multiclass Classification](#), [Classification MLPs](#)-[Classification MLPs](#)
- multidimensional scaling (MDS), [Other Dimensionality Reduction Techniques](#)
- multilabel classifiers, [Multilabel Classification](#)-[Multilabel Classification](#)
- multilayer perceptrons (MLPs), [Introduction to Artificial Neural Networks with Keras](#), [The Perceptron](#)-[Using TensorBoard for Visualization](#)
 - and autoencoders, [Efficient Data Representations](#), [Performing PCA with an Undercomplete Linear Autoencoder](#)
 - and backpropagation, [The Multilayer Perceptron and Backpropagation](#)-[The Multilayer Perceptron and Backpropagation](#)
 - callbacks, [Using Callbacks](#)-[Using Callbacks](#)
 - classification MLPs, [Classification MLPs](#)-[Classification MLPs](#)
 - complex models, [Building Complex Models Using the Functional API](#)-[Building Complex Models Using the Functional API](#)

- dynamic models, [Using the Subclassing API to Build Dynamic Models](#)-[Using the Subclassing API to Build Dynamic Models](#)
- image classifier, [Building an Image Classifier Using the Sequential API](#)-[Using the model to make predictions](#)
- regression MLPs, [Regression MLPs](#)-[Regression MLPs](#), [Building a Regression MLP Using the Sequential API](#)
- saving and restoring a model, [Saving and Restoring a Model](#)-[Saving and Restoring a Model](#)
- visualization with TensorBoard, [Using TensorBoard for Visualization](#)-[Using TensorBoard for Visualization](#)
- multimodal distribution, [Feature Scaling and Transformation](#)
- multimodal transformers, [Vision Transformers](#)
- multinomial logistic regression, [Softmax Regression](#)-[Softmax Regression](#)
- multioutput classifiers, [Multioutput Classification](#)-[Multioutput Classification](#)
- multiple regression, [Frame the Problem](#)
- multiplicative attention, [Attention Mechanisms](#)
- multitask classification, [Building Complex Models Using the Functional API](#)
- multivariate regression, [Frame the Problem](#)
- multivariate time series, [Forecasting a Time Series](#), [Forecasting Multivariate Time Series](#)-[Forecasting Multivariate Time Series](#)

N

- Nadam, [Nadam](#)
- NAG (Nesterov accelerated gradient), [Nesterov Accelerated Gradient](#), [Nadam](#)
- naive forecasting, [Forecasting a Time Series](#)
- Nash equilibrium, [The Difficulties of Training GANs](#)
- natural language processing (NLP), [Natural Language Processing with RNNs and Attention](#)-[Hugging Face's Transformers Library](#)
 - char-RNN model to generate text, [Generating Shakespearean Text Using a Character RNN](#)-[Stateful RNN](#)
 - encoder-decoder network for machine translation, [An Encoder-Decoder Network for Neural Machine Translation](#)-[An Encoder-Decoder Network for Neural Machine Translation](#)
 - machine learning examples, [Examples of Applications](#)

- sentiment analysis, [Sentiment Analysis-Reusing Pretrained Embeddings and Language Models](#)
- text classification, [Sentiment Analysis-Reusing Pretrained Embeddings and Language Models](#)
- text encoding, [Text Preprocessing-Using Pretrained Language Model Components](#)
- text summarization, [Examples of Applications](#)
- transformer models (see transformer models)
- word embeddings, [Encoding Categorical Features Using Embeddings](#)
- natural language understanding (NLU), [Examples of Applications](#)
- NCCL (Nvidia collective communications library), [Training at Scale Using the Distribution Strategies API](#)
- NEAT (neuroevolution of augmenting topologies), [Policy Search](#)
- negative class, [Confusion Matrices, Logistic Regression](#)
- nested dataset, [Preparing the Data for Machine Learning Models](#)
- Nesterov accelerated gradient (NAG), [Nesterov Accelerated Gradient, Nadam](#)
- Nesterov momentum optimization, [Nesterov Accelerated Gradient, Nadam](#)
- neural machine translation (NMT), [Reusing Pretrained Embeddings and Language Models-Multi-head attention](#)
 - and attention mechanisms, [Attention Mechanisms-Multi-head attention](#)
 - with transformers, [An Avalanche of Transformer Models-An Avalanche of Transformer Models](#)
- neural networks (see artificial neural networks)
- neuroevolution of augmenting topologies (NEAT), [Policy Search](#)
- next sentence prediction (NSP), [An Avalanche of Transformer Models](#)
- NLU (natural language understanding), [Examples of Applications](#)
- NMT (see neural machine translation)
- No Free Lunch theorem, [Data Mismatch](#)
- non-max suppression, bounding boxes, [Object Detection](#)
- nonlinear dimensionality reduction (NLDR), [LLE-LLE](#)
- nonlinear SVM classifiers, [Nonlinear SVM Classification-SVM Classes and Computational Complexity](#)
- nonparametric models, [Regularization Hyperparameters](#)
- nonrepresentative training data, [Nonrepresentative Training Data](#)
- nonresponse bias, [Nonrepresentative Training Data](#)

- normal distribution, [The Vanishing/Exploding Gradients Problems, Glorot and He Initialization](#)
- Normal equation, [The Normal Equation-The Normal Equation](#)
- normalization, [Feature Scaling and Transformation, Batch Normalization-Implementing batch normalization with Keras, Fighting the Unstable Gradients Problem](#)
- Normalization layer, [The Normalization Layer-The Normalization Layer](#)
- normalized exponential (softmax function), [Softmax Regression](#)
- notations, [Select a Performance Measure-Select a Performance Measure, Linear Regression](#)
- novelty detection, [Unsupervised learning, Using Gaussian Mixtures for Anomaly Detection](#)
- NP-Complete problem, CART as, [The CART Training Algorithm](#)
- NSP (next sentence prediction), [An Avalanche of Transformer Models](#)
- nucleus sampling, [Generating Fake Shakespearean Text](#)
- null hypothesis, [Regularization Hyperparameters](#)
- number of inputs, [Creating the model using the sequential API, Glorot and He Initialization](#)
- number of neurons per hidden layer, [Number of Neurons per Hidden Layer](#)
- NumPy, [Running the Code Examples Using Google Colab](#)
- NumPy arrays, [Clean the Data, Custom Transformers, Using TensorFlow like NumPy-Other Data Structures](#)
- Nvidia collective communications library (NCCL), [Training at Scale Using the Distribution Strategies API](#)
- Nvidia GPU card, [Getting Your Own GPU-Getting Your Own GPU](#)

O

- OAuth 2.0, [Creating a Prediction Service on Vertex AI](#)
- object detection, CNNs, [Object Detection-You Only Look Once](#)
- object tracking, CNNs, [Object Tracking](#)
- objectness score, [Object Detection](#)
- observation space, reinforcement learning, [Learning to Optimize Rewards, Neural Network Policies](#)
- observations, [Introduction to OpenAI Gym-Introduction to OpenAI Gym](#)
- OCR (optical character recognition), [The Machine Learning Landscape](#)
- OEL (open-ended learning), [Overview of Some Popular RL Algorithms](#)

- off-policy algorithm, [Q-Learning](#)
- offline learning, [Batch learning](#)
- on-policy algorithm, [Q-Learning](#)
- one-class SVM, [Other Algorithms for Anomaly and Novelty Detection](#)
- one-hot encoding, [Handling Text and Categorical Attributes](#)-[Handling Text and Categorical Attributes](#), [The CategoryEncoding Layer](#), [Encoding Categorical Features Using Embeddings](#)
- one-versus-all (OvA) strategy, [Multiclass Classification](#)-[Multiclass Classification](#)
- one-versus-one (OvO) strategy, [Multiclass Classification](#)-[Multiclass Classification](#)
- one-versus-the-rest (OvR) strategy, [Multiclass Classification](#)-[Multiclass Classification](#)
- 1cycle scheduling, [Learning Rate Scheduling](#), [Learning Rate Scheduling](#)
- 1D convolutional layers, [Semantic Segmentation](#), [Using 1D convolutional layers to process sequences](#)
- OneHotEncoder, [Handling Text and Categorical Attributes](#)-[Handling Text and Categorical Attributes](#), [Feature Scaling and Transformation](#), [Transformation Pipelines](#)
- online kernelized SVMs, [Kernelized SVMs](#)
- online learning, [Online learning](#)-[Online learning](#)
- online model, DQN, [Fixed Q-value Targets](#)
- OOB (out-of-bag) evaluation, [Out-of-Bag Evaluation](#)-[Out-of-Bag Evaluation](#)
- open-ended learning (OEL), [Overview of Some Popular RL Algorithms](#)
- OpenAI Gym, [Introduction to OpenAI Gym](#)-[Introduction to OpenAI Gym](#)
- operations (ops) and tensors, [A Quick Tour of TensorFlow](#), [Tensors and Operations](#)-[Tensors and Operations](#)
- optical character recognition (OCR), [The Machine Learning Landscape](#)
- optimal state value, MDP, [Markov Decision Processes](#)
- optimizers, [Faster Optimizers](#)-[AdamW](#)
 - AdaGrad, [AdaGrad](#)
 - Adam optimization, [Adam](#)
 - AdaMax, [AdaMax](#)
 - AdamW, [AdamW](#)
 - hyperparameters, [Learning Rate](#), [Batch Size](#), and [Other Hyperparameters](#)

- momentum optimization, [Momentum](#)
- Nadam, [Nadam](#)
- Nesterov accelerated gradient, [Nesterov Accelerated Gradient](#)
- output layer, [An Encoder–Decoder Network for Neural Machine Translation](#)
- RMSProp, [RMSProp](#)
- oracle, [Fine-Tuning Neural Network Hyperparameters](#)
- order of integration (d) hyperparameter, [The ARMA Model Family](#)
- OrdinalEncoder, [Handling Text and Categorical Attributes](#)
- orthogonal matrix, [An Encoder–Decoder Network for Neural Machine Translation](#)
- out-of-bag (OOB) evaluation, [Out-of-Bag Evaluation-Out-of-Bag Evaluation](#)
- out-of-core learning, [Online learning](#)
- out-of-sample error, [Testing and Validating](#)
- outlier detection (see anomaly detection)
- outliers, [Unsupervised Learning Techniques](#)
- output gate, LSTM, [LSTM cells](#)
- output layer, neural network, [An Encoder–Decoder Network for Neural Machine Translation](#)
- OvA (one-versus-all) strategy, [Multiclass Classification-Multiclass Classification](#)
- overcomplete autoencoder, [Convolutional Autoencoders](#)
- overfitting of data, [Overfitting the Training Data-Overfitting the Training Data, Create a Test Set](#)
 - avoiding through regularization, [Avoiding Overfitting Through Regularization-Max-Norm Regularization](#)
 - and decision trees, [Regularization Hyperparameters, Regression](#)
 - and dropout regularization, [Dropout](#)
 - gamma (γ) hyperparameter to adjust, [Gaussian RBF Kernel](#)
 - image classification, [Training and evaluating the model](#)
 - learning curves to assess, [Learning Curves-Learning Curves](#)
 - number of neurons per hidden layer, [Number of Neurons per Hidden Layer](#)
 - polynomial regression, [Training Models](#)
 - SVM model, [Soft Margin Classification](#)
- OvO (one-versus-one) strategy, [Multiclass Classification-Multiclass Classification](#)

- OvR (one-versus-the-rest) strategy, [Multiclass Classification-Multiclass Classification](#)

P

- p-value, [Regularization Hyperparameters](#)
- PACF (partial autocorrelation function), [The ARMA Model Family](#)
- padding options, convolutional layer, [Implementing Convolutional Layers with Keras](#)-[Implementing Convolutional Layers with Keras](#)
- PaLM (Pathways language model), [An Avalanche of Transformer Models](#)
- Pandas, [Running the Code Examples Using Google Colab](#), [Look for Correlations](#)-[Look for Correlations](#), [Handling Text and Categorical Attributes](#), [Feature Scaling and Transformation](#)
- parallelism, training models across devices, [Training Models Across Multiple Devices](#)-[Hyperparameter Tuning on Vertex AI](#)
 - data parallelism, [Data Parallelism-Bandwidth saturation](#)
 - distribution strategies API, [Training at Scale Using the Distribution Strategies API](#)
 - with GPU, [Parallel Execution Across Multiple Devices-Parallel Execution Across Multiple Devices](#)
 - hyperparameter tuning, [Hyperparameter Tuning on Vertex AI](#)-[Hyperparameter Tuning on Vertex AI](#)
 - model parallelism, [Model Parallelism-Model Parallelism](#)
 - on TensorFlow cluster, [Training a Model on a TensorFlow Cluster](#)-[Training a Model on a TensorFlow Cluster](#)
 - Vertex AI for running large jobs, [Running Large Training Jobs on Vertex AI](#)-[Running Large Training Jobs on Vertex AI](#)
- parameter efficiency, [Number of Hidden Layers](#)
- parameter matrix, [Softmax Regression](#)
- parameter servers, [Data parallelism with centralized parameters](#)
- parameter space, [Gradient Descent](#)
- parameter vector, [Linear Regression](#), [Gradient Descent](#), [Training and Cost Function](#), [Softmax Regression](#)
- parametric leaky ReLU (PReLU), [Leaky ReLU](#)
- parametric models, [Regularization Hyperparameters](#)
- partial autocorrelation function (PACF), [The ARMA Model Family](#)
- partial derivative, [Batch Gradient Descent](#)
- partial_fit(), [Stochastic Gradient Descent](#)

- Pathways language model (PaLM), [An Avalanche of Transformer Models](#)
- PCA (see principal component analysis)
- PDF (probability density function), [Unsupervised Learning Techniques](#), [Selecting the Number of Clusters](#)
- Pearson's r, [Look for Correlations](#)
- penalties, reinforcement learning, [Reinforcement learning](#)
- PER (prioritized experience replay), [Prioritized Experience Replay](#)
- Perceiver, [Vision Transformers](#)
- percentiles, [Take a Quick Look at the Data Structure](#)
- perceptrons, [Introduction to Artificial Neural Networks with Keras](#), [The Perceptron-Classification MLPs](#)
 - (see also multilayer perceptrons)
- performance measures, [Performance Measures-The ROC Curve](#)
 - confusion matrix, [Confusion Matrices-Confusion Matrices](#)
 - cross-validation to measure accuracy, [Measuring Accuracy Using Cross-Validation](#)-[Measuring Accuracy Using Cross-Validation](#)
 - precision and recall, [Precision and Recall-The Precision/Recall Trade-off](#)
 - ROC curve, [The ROC Curve-The ROC Curve](#)
 - selecting, [Select a Performance Measure-Select a Performance Measure](#)
- performance scheduling, [Learning Rate Scheduling](#), [Learning Rate Scheduling](#)
- permutation(), [Create a Test Set](#)
- PG (policy gradients) algorithm, [Policy Search](#), [Policy Gradients-Policy Gradients](#)
- piecewise constant scheduling, [Learning Rate Scheduling](#), [Learning Rate Scheduling](#)
- PipeDream, [Bandwidth saturation](#)
- Pipeline class, [Transformation Pipelines](#)
- Pipeline constructor, [Transformation Pipelines-Transformation Pipelines](#)
- pipelines, [Frame the Problem](#), [Transformation Pipelines-Transformation Pipelines](#), [Learning Curves](#), [Soft Margin Classification](#)
- pixelwise normalization layer, [Progressive Growing of GANs](#)
- placeholders, function definitions, [Exploring Function Definitions and Graphs](#)
- POET algorithm, [Overview of Some Popular RL Algorithms](#)

- policy gradients (PG) algorithm, [Policy Search](#), [Policy Gradients-Policy Gradients](#)
- policy parameters, [Policy Search](#)
- policy space, [Policy Search](#)
- policy, reinforcement learning, [Reinforcement learning](#), [Policy Search](#), [Neural Network Policies](#)
- polynomial features, SVM classifiers, [Polynomial Kernel](#)
- polynomial kernel, [Polynomial Kernel](#), [Kernelized SVMs](#)
- polynomial regression, [Training Models](#), [Polynomial Regression](#)-[Polynomial Regression](#)
- polynomial time, [The CART Training Algorithm](#)
- PolynomialFeatures, [Polynomial Regression](#), [Nonlinear SVM Classification](#)
- pooling kernel, [Pooling Layers](#)
- pooling layers, [Pooling Layers-Implementing Pooling Layers with Keras](#)
- positional encodings, [Positional encodings](#)-[Positional encodings](#)
- positive class, [Confusion Matrices](#), [Logistic Regression](#)
- post-training quantization, [Deploying a Model to a Mobile or Embedded Device](#)
- posterior distribution, [Variational Autoencoders](#)
- power law distribution, [Feature Scaling and Transformation](#)
- power scheduling, [Learning Rate Scheduling](#)
- PPO (proximal policy optimization), [Overview of Some Popular RL Algorithms](#)
- precision and recall, classifier metrics, [Precision and Recall-The Precision/Recall Trade-off](#)
- precision/recall curve (PR), [The Precision/Recall Trade-off](#), [The ROC Curve](#)
- precision/recall trade-off, [The Precision/Recall Trade-off-The Precision/Recall Trade-off](#)
- predict(), [Clean the Data](#), [Custom Transformers](#), [Transformation Pipelines](#)
- predicted class, [Confusion Matrices](#)
- prediction service, on Vertex AI, [Creating a Prediction Service on Vertex AI](#)-[Running Batch Prediction Jobs on Vertex AI](#)
- predictions
 - backpropagation, [The Multilayer Perceptron and Backpropagation](#)
 - confusion matrix, [Confusion Matrices](#)-[Confusion Matrices](#)

- cross-validation to measure accuracy, [Measuring Accuracy Using Cross-Validation-Measuring Accuracy Using Cross-Validation](#)
- decision trees, [Making Predictions-The CART Training Algorithm](#)
- with linear SVM classifier, [Under the Hood of Linear SVM Classifiers](#)
- predictors, [Supervised learning](#)
 - (see also ensemble learning)
- predict_log_proba(), [Soft Margin Classification](#)
- predict_proba(), [Soft Margin Classification](#)
- prefetching of data, [Prefetching-Prefetching](#)
- PReLU (parametric leaky ReLU), [Leaky ReLU](#)
- preprocessed attributes, [Take a Quick Look at the Data Structure](#)
- preprocessing data (see loading and preprocessing data)
- preprocessing mismatch, [The Normalization Layer](#)
- pretraining and pretrained layers
 - on auxiliary task, [Pretraining on an Auxiliary Task](#)
 - CNNs, [Using Pretrained Models from Keras-Pretrained Models for Transfer Learning](#)
 - greedy layer-wise pretraining, [Unsupervised Pretraining, Training One Autoencoder at a Time, Progressive Growing of GANs](#)
 - language model components, [Using Pretrained Language Model Components](#)
 - reusing embeddings, [Reusing Pretrained Embeddings and Language Models-Reusing Pretrained Embeddings and Language Models](#)
 - reusing layers, [Reusing Pretrained Layers-Pretraining on an Auxiliary Task](#)
 - in unsupervised learning, [Unsupervised Pretraining, Reusing Pretrained Embeddings and Language Models, An Avalanche of Transformer Models, Unsupervised Pretraining Using Stacked Autoencoders](#)
- primal problem, [The Dual Problem](#)
- principal component (PC), [Principal Components-Principal Components](#)
- principal component analysis (PCA), [PCA-Incremental PCA](#)
 - choosing number of dimensions, [Choosing the Right Number of Dimensions-Choosing the Right Number of Dimensions](#)
 - for compression, [PCA for Compression-PCA for Compression](#)
 - explained variance ratio, [Explained Variance Ratio](#)

- finding principal components, [Principal Components](#)
- incremental PCA, [Incremental PCA](#)-[Incremental PCA](#)
- preserving variance, [Preserving the Variance](#)
- projecting down to d dimensions, [Projecting Down to d Dimensions](#)
- randomized PCA, [Randomized PCA](#)
- for scaling data in decision trees, [Sensitivity to Axis Orientation](#)
- with undercomplete linear autoencoder, [Performing PCA with an Undercomplete Linear Autoencoder](#)-[Performing PCA with an Undercomplete Linear Autoencoder](#)
- using Scikit_Learn for, [Using Scikit-Learn](#)
- prior distribution, [Variational Autoencoders](#)
- prioritized experience replay (PER), [Prioritized Experience Replay](#)
- probabilistic autoencoders, [Variational Autoencoders](#)
- probabilities, estimating, [Estimating Probabilities](#)-[Estimating Probabilities](#), [Estimating Class Probabilities](#), [Voting Classifiers](#)
- probability density function (PDF), [Unsupervised Learning Techniques](#), [Selecting the Number of Clusters](#)
- probability versus likelihood, [Selecting the Number of Clusters](#)-[Selecting the Number of Clusters](#)
- profiling the network, with TensorBoard, [Using TensorBoard for Visualization](#)
- progressive web app (PWA), [Running a Model in a Web Page](#)
- projection, dimensionality reduction, [Projection](#)-[Projection](#)
- propositional logic, [From Biological to Artificial Neurons](#)
- protocol buffers (protobuf), [A Brief Introduction to Protocol Buffers](#)-[TensorFlow Protobufs](#), [Handling Lists of Lists Using the SequenceExample Protobuf](#), [Querying TF Serving through the gRPC API](#)
- proximal policy optimization (PPO), [Overview of Some Popular RL Algorithms](#)
- pruning of decision tree nodes, [Regularization Hyperparameters](#)
- pseudoinverse, [The Normal Equation](#)
- PWA (progressive web app), [Running a Model in a Web Page](#)
- Python API, [Get the Data](#)

Q

- Q-learning algorithm, [Q-Learning](#)-[Dueling DQN](#)
 - approximate Q-learning, [Approximate Q-Learning](#) and [Deep Q-Learning](#)

- exploration policies, [Exploration Policies](#)
- implementing deep Q-learning, [Implementing Deep Q-Learning](#)-
[Implementing Deep Q-Learning](#)
- variants in deep Q-learning, [Deep Q-Learning Variants](#)-[Dueling DQN](#)
- Q-value iteration algorithm, [Markov Decision Processes](#)-[Markov Decision Processes](#)
- Q-values, [Markov Decision Processes](#)-[Markov Decision Processes](#)
- quadratic equation, [Polynomial Regression](#)
- quadratic programming (QP) problems, [Under the Hood of Linear SVM Classifiers](#)
- quantization-aware training, [Deploying a Model to a Mobile or Embedded Device](#)
- quartiles, [Take a Quick Look at the Data Structure](#)
- queries per second (QPS), [Training and Deploying TensorFlow Models at Scale](#), [Creating a Prediction Service on Vertex AI](#)
- question-answering modules, [Examples of Applications](#)
- queues, [Other Data Structures](#), [Queues](#)

R

- radial basis function (RBF), [Feature Scaling and Transformation](#)
- ragged dimensions, [Other Data Structures](#)
- Rainbow agent, [Dueling DQN](#)
- random forests, [Better Evaluation Using Cross-Validation](#), [Ensemble Learning and Random Forests](#), [Random Forests](#)-[Feature Importance](#)
 - analysis of models and their errors, [Analyzing the Best Models and Their Errors](#)
 - decision trees (see decision trees)
 - extra-trees, [Extra-Trees](#)
 - feature importance measurement, [Feature Importance](#)
- random initialization, [Gradient Descent](#), [Gradient Descent](#)
- random patches, [Random Patches and Random Subspaces](#)
- random projection algorithm, [Random Projection](#)-[Random Projection](#)
- random subspaces, [Random Patches and Random Subspaces](#)
- RandomForestClassifier, [The ROC Curve](#)-[The ROC Curve](#)
- RandomForestRegressor, [Better Evaluation Using Cross-Validation](#)
- randomized leaky ReLU (RReLU), [Leaky ReLU](#)
- randomized PCA, [Randomized PCA](#)

- randomized search, [Randomized Search](#)-[Randomized Search](#), [Fine-Tuning Neural Network Hyperparameters](#)-[Fine-Tuning Neural Network Hyperparameters](#)
- RBF (radial basis function), [Feature Scaling and Transformation](#)
- rbf_kernel(), [Feature Scaling and Transformation](#), [Custom Transformers](#)
- recall metric, [Precision and Recall](#)
- receiver operating characteristic (ROC) curve, [The ROC Curve](#)-[The ROC Curve](#)
- recognition network, [Efficient Data Representations](#)
 - (see also autoencoders)
- recommender system, [Examples of Applications](#)
- reconstruction error, [PCA for Compression](#), [Tying Weights](#)
- reconstruction loss, [Losses and Metrics Based on Model Internals](#), [Efficient Data Representations](#)
- rectified linear units (ReLU) (see ReLU)
- recurrent dropout, [Processing Sequences Using RNNs and CNNs](#)
- recurrent layer normalization, [Processing Sequences Using RNNs and CNNs](#), [Fighting the Unstable Gradients Problem](#)
- recurrent neural networks (RNNs), [Processing Sequences Using RNNs and CNNs](#)-[Beam Search](#)
 - bidirectional, [Bidirectional RNNs](#)
 - deep RNN, [Forecasting Using a Deep RNN](#)
 - forecasting time series (see time series data)
 - gradient clipping, [Gradient Clipping](#)
 - handling long sequences, [Handling Long Sequences](#)-[WaveNet](#)
 - input and output sequences, [Input and Output Sequences](#)-[Input and Output Sequences](#)
 - memory cells, [Memory Cells](#), [Tackling the Short-Term Memory Problem](#)-[WaveNet](#)
 - NLP (see natural language processing)
 - and Perceiver, [Vision Transformers](#)
 - splitting across devices, [Model Parallelism](#)
 - stateful, [Natural Language Processing with RNNs and Attention](#), [Stateful RNN](#)-[Stateful RNN](#)
 - stateless, [Natural Language Processing with RNNs and Attention](#), [Stateful RNN](#)
 - training, [Training RNNs](#)
 - and vision transformers, [Vision Transformers](#)

- recurrent neurons, [Recurrent Neurons and Layers](#)
- reduce operation, [Data parallelism using the mirrored strategy](#)
- region proposal network (RPN), [You Only Look Once](#)
- regression MLPs, [Regression MLPs](#)-[Regression MLPs](#)
- regression models
 - and classification, [Supervised learning](#), [Multioutput Classification](#)
 - decision tree tasks, [Regression](#)-[Regression](#)
 - forecasting example, [Examples of Applications](#)
 - lasso regression, [Lasso Regression](#)-[Lasso Regression](#)
 - linear regression (see linear regression)
 - logistic regression (see logistic regression)
 - multiple regression, [Frame the Problem](#)
 - multivariate regression, [Frame the Problem](#)
 - polynomial regression, [Training Models](#), [Polynomial Regression](#)-[Polynomial Regression](#)
 - regression MLPs, [Building a Regression MLP Using the Sequential API](#)
 - ridge regression, [Ridge Regression](#)-[Ridge Regression](#), [Elastic Net Regression](#)
 - softmax regression, [Softmax Regression](#)-[Softmax Regression](#)
 - SVM, [SVM Regression](#)-[SVM Regression](#)
 - univariate regression, [Frame the Problem](#)
 - regression to the mean, [Supervised learning](#)
 - regularization, [Overfitting the Training Data](#), [Avoiding Overfitting Through Regularization](#)-[Max-Norm Regularization](#)
 - custom regularizers, [Custom Activation Functions](#), [Initializers](#), [Regularizers](#), and [Constraints](#)
 - decision trees, [Regularization Hyperparameters](#)
 - dropout, [Dropout](#)-[Dropout](#)
 - early stopping, [Early Stopping](#)-[Early Stopping](#), [Gradient Boosting](#), [Forecasting Using a Linear Model](#)
 - elastic net, [Elastic Net Regression](#)
 - hyperparameters, [Regularization Hyperparameters](#)-[Regularization Hyperparameters](#)
 - lasso regression, [Lasso Regression](#)-[Lasso Regression](#)
 - linear models, [Regularized Linear Models](#)-[Early Stopping](#)
 - max-norm, [Max-Norm Regularization](#)
 - MC dropout, [Monte Carlo \(MC\) Dropout](#)-[Monte Carlo \(MC\) Dropout](#)
 - ridge, [Ridge Regression](#)

- shrinkage, [Gradient Boosting](#)
- subword, [Sentiment Analysis](#)
- Tikhonov, [Ridge Regression-Ridge Regression](#)
- weight decay, [AdamW](#)
- ℓ_1 and ℓ_2 regularization, [\$\ell_1\$ and \$\ell_2\$ Regularization](#)
- REINFORCE algorithms, [Policy Gradients](#)
- reinforcement learning (RL), [Reinforcement learning](#), [Reinforcement Learning-Overview of Some Popular RL Algorithms](#)
 - actions, [Evaluating Actions: The Credit Assignment Problem](#)-[Evaluating Actions: The Credit Assignment Problem](#)
 - credit assignment problem, [Evaluating Actions: The Credit Assignment Problem](#)-[Evaluating Actions: The Credit Assignment Problem](#)
 - examples of, [Examples of Applications](#), [Learning to Optimize Rewards](#)
 - learning in order to optimizing rewards, [Learning to Optimize Rewards](#)
 - Markov decision processes, [Markov Decision Processes](#)-[Markov Decision Processes](#)
 - neural network policies, [Neural Network Policies](#)
 - OpenAI Gym, [Introduction to OpenAI Gym](#)-[Introduction to OpenAI Gym](#)
 - PG algorithms, [Policy Gradients](#)-[Policy Gradients](#)
 - policy search, [Policy Search](#)
 - Q-learning, [Q-Learning](#)-[Dueling DQN](#)
 - TD learning, [Temporal Difference Learning](#)
- ReLU (rectified linear units)
 - and backpropagation, [The Multilayer Perceptron and Backpropagation](#)
 - in CNN architectures, [CNN Architectures](#)
 - as default for simple tasks, [GELU, Swish, and Mish](#)
 - leakyReLU, [Leaky ReLU](#)-[Leaky ReLU](#)
 - and MLPS, [Regression MLPs](#)
 - RNN unstable gradients problem, [Fighting the Unstable Gradients Problem](#)
 - RReLU, [Leaky ReLU](#)
 - tuning hyperparameters, [Fine-Tuning Neural Network Hyperparameters](#)
- replay buffer, [Implementing Deep Q-Learning](#)

- representation learning, [Handling Text and Categorical Attributes](#)
 - (see also autoencoders)
- residual block, [Custom Models-Custom Models](#)
- residual errors, [Gradient Boosting-Gradient Boosting](#)
- residual learning, [ResNet](#)
- residual network (ResNet), [ResNet-ResNet](#)
- residual units, [ResNet](#)
- ResNet-152, [ResNet](#)
- ResNet-34, [ResNet](#), [Implementing a ResNet-34 CNN Using Keras](#)
- ResNet-50, [Using Pretrained Models from Keras](#)
- ResNeXt, [Other Noteworthy Architectures](#)
- REST API, querying through, [Querying TF Serving through the REST API](#)
- return, in reinforcement learning, [Policy Gradients](#)
- reverse process, diffusion model, [Diffusion Models-Diffusion Models](#)
- reverse-mode autodiff, [The Multilayer Perceptron and Backpropagation](#), [Computing Gradients Using Autodiff](#)
- rewards, reinforcement learning, [Reinforcement learning](#), [Learning to Optimize Rewards](#)
- Ridge, [Ridge Regression](#)
- ridge regression, [Ridge Regression-Ridge Regression](#), [Elastic Net Regression](#)
- ridge regularization, [Ridge Regression](#)
- RidgeCV, [Ridge Regression](#)
- RL (see reinforcement learning)
- RMSProp, [RMSProp](#)
- ROC (receiver operating characteristic) curve, [The ROC Curve-The ROC Curve](#)
- root mean square error (RMSE), [Select a Performance Measure-Select a Performance Measure](#), [Train and Evaluate on the Training Set](#), [Linear Regression](#), [Early Stopping](#)
- root node, decision tree, [Making Predictions](#), [Making Predictions](#)
- RPN (region proposal network), [You Only Look Once](#)
- RReLU (randomized leaky ReLU), [Leaky ReLU](#)

S

- SAC (soft actor-critic), [Overview of Some Popular RL Algorithms](#)
- “same” padding, computer vision, [Implementing Convolutional Layers with Keras](#)

- SAMME, [AdaBoost](#)
- sample inefficiency, [Policy Gradients](#)
- sampled softmax, [An Encoder–Decoder Network for Neural Machine Translation](#)
- sampling bias, [Nonrepresentative Training Data](#), [Create a Test Set](#)
- sampling noise, [Nonrepresentative Training Data](#)
- SARIMA model, [The ARMA Model Family](#)-[The ARMA Model Family](#)
- SavedModel, [Exporting SavedModels](#)-[Exporting SavedModels](#)
- saving, loading, and restoring models, [Saving and Restoring a Model](#)-[Saving and Restoring a Model](#), [Saving and Loading Models That Contain Custom Components](#)-[Saving and Loading Models That Contain Custom Components](#)
- scaled dot-product attention layer, [Multi-head attention](#)
- scatter matrix, [Look for Correlations](#)-[Look for Correlations](#)
- Scikit-Learn, [Objective and Approach](#)
 - bagging and pasting in, [Bagging and Pasting in Scikit-Learn](#)
 - CART algorithm, [The CART Training Algorithm](#), [Regression](#)
 - cross-validation, [Better Evaluation Using Cross-Validation](#)-[Better Evaluation Using Cross-Validation](#)
 - design principles, [Clean the Data](#)-[Clean the Data](#)
 - PCA implementation, [Using Scikit-Learn](#)
 - Pipeline constructor, [Transformation Pipelines](#)-[Transformation Pipelines](#)
 - sklearn.base.BaseEstimator, [Custom Transformers](#)
 - sklearn.base.clone(), [Early Stopping](#)
 - sklearn.base.TransformerMixin, [Custom Transformers](#)
 - sklearn.cluster.DBSCAN, [DBSCAN](#)
 - sklearn.cluster.KMeans, [Custom Transformers](#), [k-means](#)
 - sklearn.cluster.MiniBatchKMeans, [Accelerated k-means and mini-batch k-means](#)
 - sklearn.compose.ColumnTransformer, [Transformation Pipelines](#)
 - sklearn.compose.TransformedTargetRegressor, [Feature Scaling and Transformation](#)
 - sklearn.datasets.load_iris(), [Decision Boundaries](#)
 - sklearn.datasets.make_moons(), [Nonlinear SVM Classification](#)
 - sklearn.decomposition.IncrementalPCA, [Incremental PCA](#)
 - sklearn.decomposition.PCA, [Using Scikit-Learn](#)
 - sklearn.ensemble.AdaBoostClassifier, [AdaBoost](#)

- sklearn.ensemble.BaggingClassifier, [Bagging and Pasting in Scikit-Learn](#)
- sklearn.ensemble.GradientBoostingRegressor, [Gradient Boosting-Gradient Boosting](#)
- sklearn.ensemble.HistGradientBoostingClassifier, [Histogram-Based Gradient Boosting](#)
- sklearn.ensemble.HistGradientBoostingRegressor, [Histogram-Based Gradient Boosting](#)
- sklearn.ensemble.RandomForestClassifier, [The ROC Curve-The ROC Curve, Voting Classifiers, Random Forests, Feature Importance, Choosing the Right Number of Dimensions](#)
- sklearn.ensemble.RandomForestRegressor, [Better Evaluation Using Cross-Validation, Random Forests](#)
- sklearn.ensemble.StackingClassifier, [Stacking](#)
- sklearn.ensemble.StackingRegressor, [Stacking](#)
- sklearn.ensemble.VotingClassifier, [Voting Classifiers](#)
- sklearn.externals.joblib, [Launch, Monitor, and Maintain Your System-Launch, Monitor, and Maintain Your System](#)
- sklearn.feature_selection.SelectFromModel, [Analyzing the Best Models and Their Errors](#)
- sklearn.impute.IterativeImputer, [Clean the Data](#)
- sklearn.impute.KNNImputer, [Clean the Data](#)
- sklearn.impute.SimpleImputer, [Clean the Data](#)
- sklearn.linear_model.ElasticNet, [Elastic Net Regression](#)
- sklearn.linear_model.Lasso, [Lasso Regression](#)
- sklearn.linear_model.LinearRegression, [Model-based learning and a typical machine learning workflow, Clean the Data, Feature Scaling and Transformation, The Normal Equation, Computational Complexity, Polynomial Regression](#)
- sklearn.linear_model.LogisticRegression, [Decision Boundaries, Softmax Regression, Using Clustering for Semi-Supervised Learning](#)
- sklearn.linear_model.Perceptron, [The Perceptron](#)
- sklearn.linear_model.Ridge, [Ridge Regression](#)
- sklearn.linear_model.RidgeCV, [Ridge Regression](#)
- sklearn.linear_model.SGDClassifier, [Training a Binary Classifier, The Precision/Recall Trade-off, The Precision/Recall Trade-off, The ROC Curve-The ROC Curve, Multiclass Classification, SVM Classes and Computational Complexity, Under the Hood of Linear SVM Classifiers, The Perceptron](#)

- `sklearn.linear_model.SGDRegressor`, [Stochastic Gradient Descent](#), [Early Stopping](#)
- `sklearn.manifold.LocallyLinearEmbedding`, [LLE](#)
- `sklearn.metrics.ConfusionMatrixDisplay`, [Error Analysis](#)
- `sklearn.metrics.confusion_matrix()`, [Confusion Matrices](#), [Error Analysis](#)
- `sklearn.metrics.f1_score()`, [Precision and Recall](#), [Multilabel Classification](#)
- `sklearn.metrics.mean_squared_error()`, [Train and Evaluate on the Training Set](#)
- `sklearn.metrics.precision_recall_curve()`, [The Precision/Recall Trade-off](#), [The ROC Curve](#)
- `sklearn.metrics.precision_score()`, [Precision and Recall](#)
- `sklearn.metrics.recall_score()`, [Precision and Recall](#)
- `sklearn.metrics.roc_auc_score()`, [The ROC Curve](#)
- `sklearn.metrics.roc_curve()`, [The ROC Curve](#)
- `sklearn.metrics.silhouette_score()`, [Finding the optimal number of clusters](#)
- `sklearn.mixture.BayesianGaussianMixture`, [Bayesian Gaussian Mixture Models](#)
- `sklearn.mixture.GaussianMixture`, [Gaussian Mixtures](#)
- `sklearn.model_selection.cross_val_predict()`, [Confusion Matrices](#), [The Precision/Recall Trade-off](#), [The ROC Curve](#), [Error Analysis](#), [Stacking](#)
- `sklearn.model_selection.cross_val_score()`, [Better Evaluation Using Cross-Validation](#), [Measuring Accuracy Using Cross-Validation](#)
- `sklearn.model_selection.GridSearchCV`, [Grid Search-Grid Search](#)
- `sklearn.model_selection.learning_curve()`, [Learning Curves](#)
- `sklearn.model_selection.RandomizedSearchCV`, [Choosing the Right Number of Dimensions](#)
- `sklearn.model_selection.StratifiedKFold`, [Measuring Accuracy Using Cross-Validation](#)
- `sklearn.model_selection.StratifiedShuffleSplit`, [Create a Test Set](#)
- `sklearn.model_selection.train_test_split()`, [Create a Test Set](#), [Create a Test Set](#), [Better Evaluation Using Cross-Validation](#)
- `sklearn.multiclass.OneVsOneClassifier`, [Multiclass Classification](#)
- `sklearn.multiclass.OneVsRestClassifier`, [Multiclass Classification](#)
- `sklearn.multioutput.ChainClassifier`, [Multilabel Classification](#)

- sklearn.neighbors.KNeighborsClassifier, [Multilabel Classification](#), [Multioutput Classification](#), [DBSCAN](#)
- sklearn.neighbors.KNeighborsRegressor, [Model-based learning and a typical machine learning workflow](#)
- sklearn.neural_network.MLPClassifier, [Classification MLPs](#)
- sklearn.neural_network.MLPRegressor, [Regression MLPs](#)
- sklearn.pipeline.make_pipeline(), [Learning Curves](#)
- sklearn.pipeline.Pipeline, [Transformation Pipelines](#)
- sklearn.preprocessing.FunctionTransformer, [Custom Transformers](#)
- sklearn.preprocessing.MinMaxScaler, [Feature Scaling and Transformation](#)
- sklearn.preprocessing.OneHotEncoder, [Handling Text and Categorical Attributes](#)-[Handling Text and Categorical Attributes](#), [Feature Scaling and Transformation](#), [Transformation Pipelines](#)
- sklearn.preprocessing.OrdinalEncoder, [Handling Text and Categorical Attributes](#), [Histogram-Based Gradient Boosting](#)
- sklearn.preprocessing.PolynomialFeatures, [Polynomial Regression](#), [Nonlinear SVM Classification](#)
- sklearn.preprocessing.StandardScaler, [Feature Scaling and Transformation](#), [Gradient Descent](#), [Ridge Regression](#), [Nonlinear SVM Classification](#)
- sklearn.random_projection.GaussianRandomProjection, [Random Projection](#)
- sklearn.random_projection.SparseRandomProjection, [Random Projection](#)
- sklearn.semi_supervised.LabelPropagation, [Using Clustering for Semi-Supervised Learning](#)
- sklearn.semi_supervised.LabelSpreading, [Using Clustering for Semi-Supervised Learning](#)
- sklearn.semi_supervised.SelfTrainingClassifier, [Using Clustering for Semi-Supervised Learning](#)
- sklearn.svm.LinearSVC, [Soft Margin Classification](#), [SVM Classes and Computational Complexity](#), [Under the Hood of Linear SVM Classifiers](#)
- sklearn.svm.SVC, [Multiclass Classification](#), [Polynomial Kernel](#), [SVM Classes and Computational Complexity](#), [Under the Hood of Linear SVM Classifiers](#)
- sklearn.svm.SVR, [SVM Regression](#)

- sklearn.tree.DecisionTreeClassifier, [Training and Visualizing a Decision Tree](#), [Gini Impurity or Entropy?](#), [Regularization Hyperparameters](#), [Sensitivity to Axis Orientation](#), [Random Forests](#)
- sklearn.tree.DecisionTreeRegressor, [Train and Evaluate on the Training Set](#), [Decision Trees](#), [Regression](#), [Gradient Boosting](#)
- sklearn.tree.export_graphviz(), [Training and Visualizing a Decision Tree](#)
- sklearn.tree.ExtraTreesClassifier, [Extra-Trees](#)
- sklearn.utils.estimator_checks, [Custom Transformers](#)
- sklearn.utils.validation module, [Custom Transformers](#)
- SVM classification classes, [SVM Classes and Computational Complexity](#)
- score(), [Clean the Data](#)
- search engines, clustering for, [Clustering Algorithms: k-means and DBSCAN](#)
- search space, [Randomized Search](#)
- seasonality, time series modeling, [Forecasting a Time Series](#)
- second moment (variance of gradient), [Adam](#)
- second-order partial derivatives (Hessians), [AdamW](#)
- segment embedding, [An Avalanche of Transformer Models](#)
- SelectFromModel, [Analyzing the Best Models and Their Errors](#)
- self-attention layers, [Attention Is All You Need: The Original Transformer Architecture](#), [Multi-head attention](#), [Vision Transformers](#)
- self-distillation, [Vision Transformers](#)
- self-normalization, [ELU and SELU](#), [Dropout](#), [Summary and Practical Guidelines](#)
- self-supervised learning, [Self-supervised learning-Self-supervised learning](#)
- SELU (scaled ELU) activation function, [ELU and SELU](#), [Dropout](#)
- semantic interpolation, [Generating Fashion MNIST Images](#)
- semantic segmentation, [Examples of Applications](#), [Using Clustering for Image Segmentation](#), [Semantic Segmentation-Semantic Segmentation](#)
- semi-supervised learning, [Semi-supervised learning](#), [Clustering Algorithms: k-means and DBSCAN](#), [Using Clustering for Semi-Supervised Learning](#)-[Using Clustering for Semi-Supervised Learning](#)
- SENet, [SENet-SENet](#)
- sensitivity (recall), ROC curve, [The ROC Curve](#)
- sensitivity metric, [Confusion Matrices](#)
- sensors, [Learning to Optimize Rewards](#)

- sentence encoder, [Using Pretrained Language Model Components](#), [Reusing Pretrained Embeddings and Language Models](#)
- SentencePiece project, [Sentiment Analysis](#)
- sentiment analysis, [Sentiment Analysis-Reusing Pretrained Embeddings and Language Models](#)
- sentiment neuron, [Stateful RNN](#)
- separable convolution layer, [Xception](#)
- sequence length, [Using 1D convolutional layers to process sequences](#), [Positional encodings](#)
- sequence-to-sequence (seq2seq) network, [Input and Output Sequences](#), [Forecasting Using a Sequence-to-Sequence Model](#)-[Forecasting Using a Sequence-to-Sequence Model](#)
- sequence-to-vector network, [Input and Output Sequences](#)
- SequenceExample protobuf, [Handling Lists of Lists Using the SequenceExample Protobuf](#)
- sequential API, image classifier with, [Building an Image Classifier Using the Sequential API](#)-[Using the model to make predictions](#)
- service account, GCP, [Creating a Prediction Service on Vertex AI](#)
- service worker, [Running a Model in a Web Page](#)
- sets, [Sets](#)
- set_params(), [Custom Transformers](#)
- SGD (see stochastic gradient descent)
- SGDClassifier, [Training a Binary Classifier](#), [The Precision/Recall Trade-off](#), [The Precision/Recall Trade-off](#), [The ROC Curve](#)-[The ROC Curve](#), [Multiclass Classification](#), [SVM Classes and Computational Complexity](#), [Under the Hood of Linear SVM Classifiers](#), [Under the Hood of Linear SVM Classifiers](#)
- SGDRegressor, [Stochastic Gradient Descent](#), [Early Stopping](#)
- sharpening, NLP transformers, [Vision Transformers](#)
- shrinkage, [Gradient Boosting](#)
- shuffle_and_split_data(), [Create a Test Set](#), [Create a Test Set](#)
- shuffling data, [Stochastic Gradient Descent](#), [Shuffling the Data](#)-[Shuffling the Data](#)
- Siamese neural network, [Parallel Execution Across Multiple Devices](#)
- sigmoid activation function, [Estimating Probabilities](#), [The Multilayer Perceptron and Backpropagation](#), [The Vanishing/Exploding Gradients Problems](#), [Sparse Autoencoders](#)
- signals, [Biological Neurons](#)
- silhouette coefficient, [Finding the optimal number of clusters](#)

- silhouette diagram, [Finding the optimal number of clusters](#)
- SiLU activation function, [GELU, Swish, and Mish](#)
- similarity features, SVM, [Similarity Features](#)
- SimpleImputer, [Clean the Data](#)
- simulated annealing, [Stochastic Gradient Descent](#)
- simulated environments, [Introduction to OpenAI Gym](#)-[Introduction to OpenAI Gym](#)
- single program, multiple data (SPMD), [Data Parallelism](#)-[Bandwidth saturation](#)
- single-shot learning, [Semantic Segmentation](#)
- singular value decomposition (SVD), [The Normal Equation](#), [Principal Components](#), [Randomized PCA](#)
- skewed datasets, [Measuring Accuracy Using Cross-Validation](#)
- skewed left/right, [Take a Quick Look at the Data Structure](#)
- skip connections, [ELU and SELU](#), [ResNet](#)
- slack variable, [Under the Hood of Linear SVM Classifiers](#)
- smoothing terms, [Batch Normalization](#), [AdaGrad](#)
- SMOTE (synthetic minority oversampling technique), [AlexNet](#)
- soft actor-critic (SAC), [Overview of Some Popular RL Algorithms](#)
- soft clustering, [k-means](#)
- soft margin classification, [Soft Margin Classification](#)-[Soft Margin Classification](#), [Under the Hood of Linear SVM Classifiers](#)
- soft voting, [Voting Classifiers](#)
- softmax activation function, [Softmax Regression](#), [Classification MLPs](#), [Creating the model using the sequential API](#), [An Encoder–Decoder Network for Neural Machine Translation](#)
- softmax regression, [Softmax Regression](#)-[Softmax Regression](#)
- softplus activation function, [Regression MLPs](#)
- spam filters, [The Machine Learning Landscape](#)-[Why Use Machine Learning?](#), [Types of Machine Learning Systems](#)-[Supervised learning](#)
- sparse autoencoders, [Sparse Autoencoders](#)-[Sparse Autoencoders](#)
- sparse features, SVC class, [SVM Classes and Computational Complexity](#)
- sparse matrix, [Handling Text and Categorical Attributes](#), [Feature Scaling and Transformation](#), [Transformation Pipelines](#)
- sparse models, [Lasso Regression](#), [AdamW](#)
- sparse tensors, [Sparse Tensors](#)
- sparsity loss, [Sparse Autoencoders](#)
- specificity, ROC curve, [The ROC Curve](#)
- spectral clustering, [Other Clustering Algorithms](#)

- speech recognition, [Why Use Machine Learning? Examples of Applications](#)
- split node, decision tree, [Making Predictions](#)
- SPMD (single program, multiple data), [Data Parallelism-Bandwidth saturation](#)
- squared hinge loss, [Under the Hood of Linear SVM Classifiers](#)
- Stable Diffusion, [Diffusion Models](#)
- stacked autoencoders, [Stacked Autoencoders-Training One Autoencoder at a Time](#)
- stacking (stacked generalization), [Stacking-Stacking](#)
- stale gradients, [Asynchronous updates](#)
- standard correlation coefficient, [Look for Correlations](#)
- standard deviation, [Take a Quick Look at the Data Structure](#)
- standardization, [Feature Scaling and Transformation](#)
- StandardScaler, [Feature Scaling and Transformation](#), [Gradient Descent](#), [Ridge Regression](#), [Nonlinear SVM Classification](#)
- state-action values (Q-Values), [Markov Decision Processes](#)-[Markov Decision Processes](#)
- stateful metric, [Custom Metrics](#)
- stateful RNN, [Natural Language Processing with RNNs and Attention](#), [Stateful RNN-Stateful RNN](#)
- stateless RNN, [Natural Language Processing with RNNs and Attention](#), [Stateful RNN](#), [Stateful RNN](#)
- stationary time series, [Forecasting a Time Series](#)
- statistical mode, [Bagging and Pasting](#)
- statistical significance, [Regularization Hyperparameters](#)
- statsmodels library, [The ARMA Model Family](#)
- stemming, [Exercises](#)
- step functions, TLU, [The Perceptron](#)
- stochastic gradient boosting, [Gradient Boosting](#)
- stochastic gradient descent (SGD), [Stochastic Gradient Descent](#)-[Stochastic Gradient Descent](#), [SVM Classes and Computational Complexity](#)
 - early stopping, [Early Stopping](#)
 - image classification, [Compiling the model](#)
 - and perceptron learning algorithm, [The Perceptron](#)
 - ridge regularization, [Ridge Regression](#)
 - and TD learning, [Temporal Difference Learning](#)
 - training binary classifier, [Training a Binary Classifier](#)

- stratified sampling, [Create a Test Set](#)-[Create a Test Set](#), [Measuring Accuracy Using Cross-Validation](#)
- strat_train_set, [Prepare the Data for Machine Learning Algorithms](#)
- streaming metric, [Custom Metrics](#)
- strides, [Convolutional Layers](#), [Semantic Segmentation](#), [Using 1D convolutional layers to process sequences](#)
- string kernels, [Gaussian RBF Kernel](#)
- string tensors, [Other Data Structures](#)
- StringLookup layer, [The StringLookup Layer](#)
- strings, [Strings](#)-[Strings](#)
- strong learners, [Voting Classifiers](#)
- style mixing, StyleGAN, [StyleGANs](#)
- style transfer, GANs, [StyleGANs](#)
- StyleGANs, [StyleGANs](#)-[StyleGANs](#)
- subclassing API, [Using the Subclassing API to Build Dynamic Models](#)-[Using the Subclassing API to Build Dynamic Models](#)
- subgradient vector, [Lasso Regression](#)
- subsampling, pooling layer, [Pooling Layers](#)
- subword regularization, [Sentiment Analysis](#)
- super-convergence, [Learning Rate Scheduling](#)
- super-resolution, [Semantic Segmentation](#)
- supervised learning, [Supervised learning](#)
- support vector machines (SVMs), [Support Vector Machines](#)-[Kernelized SVMs](#)
 - decision function, [Under the Hood of Linear SVM Classifiers](#)-[Under the Hood of Linear SVM Classifiers](#)
 - dual problem, [The Dual Problem](#)-[Kernelized SVMs](#)
 - kernelized SVMs, [Kernelized SVMs](#)-[Kernelized SVMs](#)
 - linear classification, [Linear SVM Classification](#)-[Soft Margin Classification](#)
 - mechanics of, [Under the Hood of Linear SVM Classifiers](#)-[Kernelized SVMs](#)
 - in multiclass classification, [Multiclass Classification](#)
 - nonlinear classifiers, [Nonlinear SVM Classification](#)
 - one-class SVM, [Other Algorithms for Anomaly and Novelty Detection](#)
 - SVM regression, [SVM Regression](#)-[SVM Regression](#)
- support vectors, [Linear SVM Classification](#)

- SVC class, [Polynomial Kernel](#), [SVM Classes and Computational Complexity](#)
- SVD (singular value decomposition), [The Normal Equation](#), [Principal Components](#), [Randomized PCA](#)
- SVMs (see support vector machines)
- SVR class, [SVM Regression](#)
- Swish activation function, [GELU, Swish, and Mish](#)
- Swiss roll dataset, [Manifold Learning-Manifold Learning](#)
- symbolic differentiation, [Forward-Mode Autodiff](#)
- symbolic tensors, [AutoGraph and Tracing](#), [TF Functions and Concrete Functions](#)
- synchronous updates, with centralized parameters, [Synchronous updates](#)
- synthetic minority oversampling technique (SMOTE), [AlexNet](#)

T

- t-distributed stochastic neighbor embedding (t-SNE), [Other Dimensionality Reduction Techniques](#), [Visualizing the Fashion MNIST Dataset](#)
- target attributes, [Take a Quick Look at the Data Structure](#)
- target distribution, transforming, [Feature Scaling and Transformation](#)
- target model, DQN, [Fixed Q-value Targets](#)
- TD error, [Temporal Difference Learning](#)
- TD target, [Temporal Difference Learning](#)
- TD-Gammon, [Reinforcement Learning](#)
- teacher forcing, [An Encoder–Decoder Network for Neural Machine Translation](#)
- temperature, Char-RNN model, [Generating Fake Shakespearean Text](#)
- temporal difference (TD) learning, [Temporal Difference Learning](#), [Prioritized Experience Replay](#)
- tensor arrays, [Tensor Arrays](#)
- tensor processing units (TPUs), [A Quick Tour of TensorFlow](#), [Deploying a new model version](#), [Deploying a Model to a Mobile or Embedded Device](#), [Training a Model on a TensorFlow Cluster](#)
- TensorBoard, [Using TensorBoard for Visualization](#)-[Using TensorBoard for Visualization](#), [A Quick Tour of TensorFlow](#), [Masking](#), [Running Large Training Jobs on Vertex AI](#)
- TensorFlow, [Objective and Approach](#), [Objective and Approach](#), [Custom Models and Training with TensorFlow](#)-[The TensorFlow Datasets](#)

[Project, Training and Deploying TensorFlow Models at Scale](#)

[Hyperparameter Tuning on Vertex AI](#)

- (see also Keras API)
- architecture, [A Quick Tour of TensorFlow](#)
- creating training function, [Using the Dataset with Keras](#)
- custom models (see custom models and training algorithms)
- deploying model to a mobile device, [Deploying a Model to a Mobile or Embedded Device](#)-[Deploying a Model to a Mobile or Embedded Device](#)
- functions and graphs, [TensorFlow Graphs-Using TF Functions with Keras \(or Not\)](#)
- GPU management with, [Managing the GPU RAM](#)-[Managing the GPU RAM](#), [Parallel Execution Across Multiple Devices](#)-[Parallel Execution Across Multiple Devices](#)
- graphs and functions, [A Quick Tour of TensorFlow](#), [TensorFlow Functions and Graphs](#)-[TF Function Rules](#), [TensorFlow Graphs](#)-[Using TF Functions with Keras \(or Not\)](#)
- hub.KerasLayer, [Using Pretrained Language Model Components](#)
- math operations, [Tensors and Operations](#)
- with NumPy, [Using TensorFlow like NumPy](#)-[Other Data Structures](#)
- operations (ops) and tensors, [A Quick Tour of TensorFlow](#), [Tensors and Operations](#)-[Tensors and NumPy](#)
- parallelism to train models (see parallelism)
- platforms and APIs available, [A Quick Tour of TensorFlow](#)
- serving a model (see TensorFlow Serving)
- special data structures, [Special Data Structures](#)-[Queues](#)
- tf.add(), [Tensors and Operations](#)
- tf.autograph.to_code(), [AutoGraph and Tracing](#)
- tf.cast(), [Type Conversions](#)
- tf.config.set_soft_device_placement, [Placing Operations and Variables on Devices](#)
- tf.config.threading.set_inter_op_parallelism_threads(), [Parallel Execution Across Multiple Devices](#)
- tf.config.threading.set_intra_op_parallelism_threads(), [Parallel Execution Across Multiple Devices](#)
- tf.constant(), [Tensors and Operations](#)
- tf.data API (see tf.data API)
- tf.device(), [Placing Operations and Variables on Devices](#)

- `tf.distribute.experimental.CentralStorageStrategy`, [Training at Scale Using the Distribution Strategies API](#)
- `tf.distribute.experimental.TPUStrategy`, [Training a Model on a TensorFlow Cluster](#)
- `tf.distribute.MirroredStrategy`, [Training at Scale Using the Distribution Strategies API](#), [Hyperparameter Tuning on Vertex AI](#)
- `tf.distribute.MultiWorkerMirroredStrategy`, [Training a Model on a TensorFlow Cluster](#)
- `tf.float32`, [Tensors and NumPy](#)
- `tf.float32` data type, [Custom Metrics](#), [Preprocessing the Data](#)
- `tf.function()`, [TensorFlow Functions and Graphs](#), [TF Function Rules](#)
- `tf.int32` data type, [Other Data Structures](#)
- `tf.io.decode_base64()`, [Running Batch Prediction Jobs on Vertex AI](#)
- `tf.io.decode_csv()`, [Preprocessing the Data](#)
- `tf.io.decode_image()`, [Running Batch Prediction Jobs on Vertex AI](#)
- `tf.io.decode_png()`, [Running Batch Prediction Jobs on Vertex AI](#)
- `tf.io.decode_proto()`, [A Brief Introduction to Protocol Buffers](#)
- `tf.io.FixedLenFeature`, [Loading and Parsing Examples](#)
- `tf.io.parse_example()`, [Loading and Parsing Examples](#)
- `tf.io.parse_sequence_example()`, [Handling Lists of Lists Using the SequenceExample Protobuf](#)
- `tf.io.parse_single_example()`, [Loading and Parsing Examples](#)
- `tf.io.parse_single_sequence_example()`, [Handling Lists of Lists Using the SequenceExample Protobuf](#)
- `tf.io.parse_tensor()`, [Loading and Parsing Examples](#)
- `tf.io.serialize_tensor()`, [Loading and Parsing Examples](#)
- `tf.io.TFRecordOptions`, [Compressed TFRecord Files](#)
- `tf.io.TFRecordWriter`, [The TFRecord Format](#)
- `tf.io.VarLenFeature`, [Loading and Parsing Examples](#)
- `tf.linalg.band_part()`, [Multi-head attention](#)
- `tf.lite.TFLiteConverter.from_keras_model()`, [Deploying a Model to a Mobile or Embedded Device](#)
- `tf.make_tensor_proto()`, [Querying TF Serving through the gRPC API](#)
- `tf.matmul()`, [Tensors and Operations](#), [Multi-head attention](#)
- `tf.nn.conv2d()`, [Implementing Convolutional Layers with Keras](#)
- `tf.nn.embedding_lookup()`, [Image Preprocessing Layers](#)
- `tf.nn.local_response_normalization()`, [AlexNet](#)
- `tf.nn.moments()`, [Image Preprocessing Layers](#)

- tf.nn.sampled_softmax_loss(), [An Encoder–Decoder Network for Neural Machine Translation](#)
- tf.py_function(), [TF Function Rules](#), [Exporting SavedModels](#)
- tf.queue module, [Other Data Structures](#)
- tf.queue.FIFOQueue, [Queues](#)
- tf.RaggedTensor, [Other Data Structures](#)
- tf.random.categorical(), [Generating Fake Shakespearean Text](#)
- tf.reduce_max(), [Implementing Pooling Layers with Keras](#)
- tf.reduce_mean(), [Custom Training Loops](#)
- tf.reduce_sum(), [TensorFlow Functions and Graphs](#)
- tf.saved_model_cli command, [Exporting SavedModels](#)
- tf.sets module, [Other Data Structures](#)
- tf.sort(), [TF Function Rules](#)
- tf.SparseTensor, [Other Data Structures](#)
- tf.stack(), [Preprocessing the Data](#), [Stateful RNN](#)
- tf.string data type, [Other Data Structures](#)
- tf.strings module, [Other Data Structures](#)
- tf.Tensor, [Tensors and Operations](#), [Variables](#)
- tf.TensorArray, [Other Data Structures](#)
- tf.transpose(), [Tensors and Operations](#)
- tf.Variable, [Variables](#)
- tf.Variable.assign(), [Variables](#)
- type conversions, [Type Conversions](#)
- variables, [Variables](#)
- web page, running a model in, [Running a Model in a Web Page](#)
- TensorFlow cluster, [Training a Model on a TensorFlow Cluster](#)-
[Training a Model on a TensorFlow Cluster](#)
- TensorFlow Datasets (TFDS) project, [The TensorFlow Datasets Project](#)-
[The TensorFlow Datasets Project](#)
- TensorFlow Extended (TFX), [A Quick Tour of TensorFlow](#)
- TensorFlow Hub, [A Quick Tour of TensorFlow](#), [Using Pretrained Language Model Components](#), [Reusing Pretrained Embeddings and Language Models](#)
- TensorFlow Lite, [A Quick Tour of TensorFlow](#)
- TensorFlow playground, [Classification MLPs](#)
- TensorFlow Serving (TF Serving), [Serving a TensorFlow Model](#)-
[Running Batch Prediction Jobs on Vertex AI](#)
 - batch prediction jobs on Vertex AI, [Running Batch Prediction Jobs on Vertex AI](#)

- creating prediction service, [Creating a Prediction Service on Vertex AI](#)-
[Creating a Prediction Service on Vertex AI](#)
- deploying new model version, [Deploying a new model version](#)-
[Deploying a new model version](#)
- Docker container, [Installing and starting TensorFlow Serving](#)
- exporting SavedModels, [Exporting SavedModels](#)-[Exporting SavedModels](#)
- gRPC API, querying through, [Querying TF Serving through the gRPC API](#)
- installing and starting up, [Installing and starting TensorFlow Serving](#)-[Installing and starting TensorFlow Serving](#)
- REST API, querying through, [Querying TF Serving through the REST API](#)
- TensorFlow Text, [Text Preprocessing](#), [Sentiment Analysis](#)
- TensorFlow.js (TFJS) JavaScript library, [Running a Model in a Web Page](#)
- tensors, [Tensors and Operations](#)-[Tensors and NumPy](#)
- term-frequency x inverse-document-frequency (TF-IDF), [Text Preprocessing](#)
- terminal state, Markov chain, [Markov Decision Processes](#)
- test set, [Testing and Validating](#), [Create a Test Set](#)-[Create a Test Set](#)
- text attributes, [Handling Text and Categorical Attributes](#)
- text processing (see natural language processing)
- TF Serving (see TensorFlow Serving)
- tf.data API, [The tf.data API](#)-[Using the Dataset with Keras](#)
 - chaining transformations, [Chaining Transformations](#)-[Chaining Transformations](#)
 - interleaving lines from multiple files, [Interleaving Lines from Multiple Files](#)-[Interleaving Lines from Multiple Files](#)
 - and Keras preprocessing layers, [The Normalization Layer](#)
 - prefetching, [Prefetching](#)-[Prefetching](#)
 - preprocessing the data, [Preprocessing the Data](#)-[Preprocessing the Data](#)
 - shuffling data, [Shuffling the Data](#)-[Shuffling the Data](#)
 - tf.data.AUTOTUNE, [Chaining Transformations](#)
 - tf.data.Dataset.from_tensor_slices(), [The tf.data API](#)-[Chaining Transformations](#)
 - tf.data.TFRecordDataset, [The TFRecord Format](#), [The TFRecord Format](#), [Loading and Parsing Examples](#)

- using dataset with Keras, [Using the Dataset with Keras](#)-[Using the Dataset with Keras](#)
- TFDS (TensorFlow Datasets) project, [The TensorFlow Datasets Project](#)-[The TensorFlow Datasets Project](#)
- TFJS (TensorFlow.js) JavaScript library, [Running a Model in a Web Page](#)
- TFLite, [Deploying a Model to a Mobile or Embedded Device](#)-[Deploying a Model to a Mobile or Embedded Device](#)
- TFRecord format, [Loading and Preprocessing Data with TensorFlow](#), [The TFRecord Format](#)-[Handling Lists of Lists Using the SequenceExample Protobuf](#)
- TFX (TensorFlow Extended), [A Quick Tour of TensorFlow](#)
- theoretical information criterion, [Selecting the Number of Clusters](#)
- 3D convolutional layers, [Semantic Segmentation](#)
- threshold logic units (TLUs), [The Perceptron](#), [The Multilayer Perceptron and Backpropagation](#)
- Tikhonov regularization, [Ridge Regression](#)-[Ridge Regression](#), [Elastic Net Regression](#)
- time series data, forecasting, [Processing Sequences Using RNNs and CNNs](#), [Forecasting a Time Series](#)-[Forecasting Using a Sequence-to-Sequence Model](#)
 - ARMA model family, [The ARMA Model Family](#)-[The ARMA Model Family](#)
 - data preparation for ML models, [Preparing the Data for Machine Learning Models](#)-[Preparing the Data for Machine Learning Models](#)
 - with deep RNN, [Forecasting Using a Deep RNN](#)
 - with linear model, [Forecasting Using a Linear Model](#)
 - multivariate time series, [Forecasting a Time Series](#), [Forecasting Multivariate Time Series](#)-[Forecasting Multivariate Time Series](#)
 - with sequence-to-sequence model, [Input and Output Sequences](#), [Forecasting Using a Sequence-to-Sequence Model](#)-[Forecasting Using a Sequence-to-Sequence Model](#)
 - several time steps ahead, [Forecasting Several Time Steps Ahead](#)-[Forecasting Several Time Steps Ahead](#)
 - with simple RNN, [Forecasting Using a Simple RNN](#)-[Forecasting Using a Simple RNN](#)
- TLUs (threshold logic units), [The Perceptron](#), [The Multilayer Perceptron and Backpropagation](#)
- TNR (true negative rate), [The ROC Curve](#)

- tokenizers library, [Sentiment Analysis](#)
- tolerance (ϵ), [Batch Gradient Descent](#), [SVM Classes and Computational Complexity](#)
- TPR (true positive rate), [Confusion Matrices](#), [The ROC Curve](#)
- TPUs (tensor processing units), [A Quick Tour of TensorFlow](#), [Deploying a new model version](#), [Deploying a Model to a Mobile or Embedded Device](#), [Training a Model on a TensorFlow Cluster](#)
- train-dev set, [Data Mismatch](#)
- training, [Model-based learning and a typical machine learning workflow](#)
- training instance, [What Is Machine Learning?](#)
- training loops, [Custom Training Loops-Custom Training Loops](#), [Using the Dataset with Keras](#)
- training models, [Training Models-Softmax Regression](#)
 - learning curves in, [Learning Curves-Learning Curves](#)
 - linear regression, [Training Models](#), [Linear Regression-Mini-Batch Gradient Descent](#)
 - logistic regression, [Logistic Regression-Softmax Regression](#)
 - perceptrons, [The Perceptron-The Perceptron](#)
 - polynomial regression, [Training Models](#), [Polynomial Regression-Polynomial Regression](#)
- training set, [What Is Machine Learning?](#), [Testing and Validating](#)
 - cost function of, [Training and Cost Function-Training and Cost Function](#)
 - insufficient quantities, [Insufficient Quantity of Training Data](#)
 - irrelevant features, [Irrelevant Features](#)
 - min-max scaling, [Feature Scaling and Transformation](#)
 - nonrepresentative, [Nonrepresentative Training Data](#)
 - overfitting, [Overfitting the Training Data-Overfitting the Training Data](#)
 - preparing for ML algorithms, [Prepare the Data for Machine Learning Algorithms](#)
 - training and evaluating on, [Train and Evaluate on the Training Set-Train and Evaluate on the Training Set](#)
 - transforming data, [Feature Scaling and Transformation](#)
 - underfitting, [Underfitting the Training Data](#)
 - visualizing data, [Explore and Visualize the Data to Gain Insights](#)
- training set expansion, [Exercises](#), [AlexNet](#), [Pretrained Models for Transfer Learning](#)

- training/serving skew, [Loading and Preprocessing Data with TensorFlow](#)
- train_test_split(), [Create a Test Set](#), [Better Evaluation Using Cross-Validation](#)
- transfer learning, [Self-supervised learning](#), [Reusing Pretrained Layers](#), [Transfer Learning with Keras](#)-[Transfer Learning with Keras](#), [Pretrained Models for Transfer Learning](#)-[Pretrained Models for Transfer Learning](#)
- transform(), [Clean the Data](#), [Handling Text and Categorical Attributes](#), [Feature Scaling and Transformation](#), [Custom Transformers](#)
- transformation of data
 - custom transformers, [Custom Transformers](#)-[Custom Transformers](#)
 - estimator transformers, [Clean the Data](#)
 - and feature scaling, [Feature Scaling and Transformation](#)-[Feature Scaling and Transformation](#)
 - transformer models (see transformer models)
- transformation pipelines, [Transformation Pipelines](#)-[Transformation Pipelines](#)
- TransformedTargetRegressor, [Feature Scaling and Transformation](#)
- transformer, [Attention Is All You Need: The Original Transformer Architecture](#)
- transformer models
 - attention mechanisms, [Attention Is All You Need: The Original Transformer Architecture](#)-[Multi-head attention](#), [Vision Transformers](#)
 - BERT, [An Avalanche of Transformer Models](#)
 - DistilBERT, [An Avalanche of Transformer Models](#), [Hugging Face's Transformers Library](#)-[Hugging Face's Transformers Library](#)
 - Hugging Face library, [Hugging Face's Transformers Library](#)-[Hugging Face's Transformers Library](#)
 - Pathways language model, [An Avalanche of Transformer Models](#)
 - vision transformers, [Vision Transformers](#)-[Vision Transformers](#)
- TransformerMixin, [Custom Transformers](#)
- transformers library, [Hugging Face's Transformers Library](#)-[Hugging Face's Transformers Library](#)
- translation, with RNNs, [Natural Language Processing with RNNs and Attention](#), [An Encoder–Decoder Network for Neural Machine Translation](#)-[Beam Search](#)
 - (see also transformer models)

- transpose operator, [Select a Performance Measure](#)
- transposed convolutional layer, [Semantic Segmentation-Semantic Segmentation](#)
- true negative rate (TNR), [The ROC Curve](#)
- true negatives, confusion matrix, [Confusion Matrices](#)
- true positive rate (TPR), [Confusion Matrices](#), [The ROC Curve](#)
- true positives, confusion matrix, [Confusion Matrices](#)
- trust region policy optimization (TRPO), [Overview of Some Popular RL Algorithms](#)
- 2D convolutional layers, [Implementing Convolutional Layers with Keras](#)
- tying weights, [Tying Weights](#)
- type I errors, confusion matrix, [Confusion Matrices](#)
- type II errors, confusion matrix, [Confusion Matrices](#)

U

- uncertainty sampling, [Using Clustering for Semi-Supervised Learning](#)
- undercomplete, autoencoder as, [Efficient Data Representations-Performing PCA with an Undercomplete Linear Autoencoder](#)
- underfitting of data, [Underfitting the Training Data](#), [Train and Evaluate on the Training Set](#), [Learning Curves-Learning Curves](#), [Gaussian RBF Kernel](#)
- univariate regression, [Frame the Problem](#)
- univariate time series, [Forecasting a Time Series](#)
- Universal Sentence Encoder, [Reusing Pretrained Embeddings and Language Models-Reusing Pretrained Embeddings and Language Models](#)
- unreasonable effectiveness of data, [Insufficient Quantity of Training Data](#)
- unrolling the network through time, [Recurrent Neurons and Layers](#)
- unstable gradients problem, [The Vanishing/Exploding Gradients Problems, Fighting the Unstable Gradients Problem](#)
 - (see also vanishing and exploding gradients)
- unsupervised learning, [Unsupervised learning-Unsupervised learning](#), [Unsupervised Learning Techniques-Other Algorithms for Anomaly and Novelty Detection](#)
 - anomaly detection, [Unsupervised learning](#)
 - association rule learning, [Unsupervised learning](#)
 - autoencoders (see autoencoders)

- clustering (see clustering algorithms)
- density estimation, [Unsupervised Learning Techniques](#)
- diffusion models, [Diffusion Models-Diffusion Models](#)
- dimensionality reduction (see dimensionality reduction)
- GANs (see generative adversarial networks)
- GMM, [Gaussian Mixtures-Other Algorithms for Anomaly and Novelty Detection](#)
- k-means (see k-means algorithm)
- novelty detection, [Unsupervised learning](#)
- pretraining, [Unsupervised Pretraining, Reusing Pretrained Embeddings and Language Models](#), [An Avalanche of Transformer Models](#), [Unsupervised Pretraining Using Stacked Autoencoders](#)
- stacked autoencoders, [Unsupervised Pretraining Using Stacked Autoencoders](#)
- transformer models, [An Avalanche of Transformer Models](#)
- visualization algorithms, [Unsupervised learning-Unsupervised learning](#)
- upsampling layer, [Semantic Segmentation](#)
- utility function, [Model-based learning and a typical machine learning workflow](#)

V

- VAEs (variational autoencoders), [Variational Autoencoders-Variational Autoencoders](#)
- “valid” padding, computer vision, [Implementing Convolutional Layers with Keras](#)
- validation set, [Hyperparameter Tuning and Model Selection](#), [Training and evaluating the model](#)-[Training and evaluating the model](#)
- value_counts(), [Take a Quick Look at the Data Structure](#)
- vanishing and exploding gradients, [The Vanishing/Exploding Gradients Problems-Gradient Clipping](#)
 - activation function improvements, [Better Activation Functions-GELU, Swish, and Mish](#)
 - batch normalization, [Batch Normalization-Implementing batch normalization with Keras](#)
 - Glorot and He initialization, [Glorot and He Initialization-Glorot and He Initialization](#)
 - gradient clipping, [Gradient Clipping](#)

- unstable gradients problem, [Fighting the Unstable Gradients Problem](#)
- variables
 - handling in TF functions, [Handling Variables and Other Resources in TF Functions](#)
 - persistence of, [Custom Metrics](#)
 - placing on GPUs, [Placing Operations and Variables on Devices](#)
 - in TensorFlow, [Variables](#)
- variance
 - bias/variance trade-off, [Learning Curves](#)
 - explained, [Explained Variance Ratio-Choosing the Right Number of Dimensions](#)
 - high variance with decision trees, [Decision Trees Have a High Variance](#)
 - preserving, [Preserving the Variance](#)
- variational autoencoders (VAEs), [Variational Autoencoders-Variational Autoencoders](#)
- vector-to-sequence network, [Input and Output Sequences](#)
- vectors, norms for measuring distance, [Select a Performance Measure](#)
- Vertex AI, [Launch, Monitor, and Maintain Your System](#), [Creating a Prediction Service on Vertex AI-Creating a Prediction Service on Vertex AI](#), [Running Large Training Jobs on Vertex AI-Running Large Training Jobs on Vertex AI](#)
- VGGNet, [VGGNet](#)
- virtual GPU device, [Managing the GPU RAM](#)
- vision transformers (ViTs), [Vision Transformers-Vision Transformers](#)
- visual cortex architecture, [The Architecture of the Visual Cortex](#)
- visualization of data, [Examples of Applications](#), [Unsupervised learning-Unsupervised learning](#), [Explore and Visualize the Data to Gain Insights-Experiment with Attribute Combinations](#)
 - decision trees, [Training and Visualizing a Decision Tree-Training and Visualizing a Decision Tree](#)
 - dimensionality reduction, [Dimensionality Reduction, Choosing the Right Number of Dimensions](#)
 - end-to-end exercise, [Explore and Visualize the Data to Gain Insights-Experiment with Attribute Combinations](#)
 - MLPs with TensorBoard, [Using TensorBoard for Visualization-Using TensorBoard for Visualization](#)

- stacked autoencoders, [Visualizing the Reconstructions](#)-[Visualizing the Reconstructions](#)
- t-SNE, [Other Dimensionality Reduction Techniques](#)
- ViTs (vision transformers), [Vision Transformers](#)-[Vision Transformers](#)
- voting classifiers, [Voting Classifiers](#)-[Voting Classifiers](#)

W

- wall time, [Batch Normalization](#)
- warmup phase, asynchronous model updates, [Asynchronous updates](#)
- WaveNet, [Processing Sequences Using RNNs and CNNs](#), [WaveNet](#)-[WaveNet](#)
- weak learners, [Voting Classifiers](#)
- web page, running a model in, [Running a Model in a Web Page](#)
- weight decay, [AdamW](#)
- weight stashing, [Bandwidth saturation](#)
- weight-tying, [Tying Weights](#)
- weights
 - boosting, [AdaBoost](#)-[AdaBoost](#)
 - convolutional layers, [Implementing Convolutional Layers with Keras](#)
 - freezing reused layers, [Reusing Pretrained Layers](#)
 - of hidden layers, [Creating the model using the sequential API](#)
 - in prioritized experience replay, [Prioritized Experience Replay](#)
 - saving instead of whole model, [Saving and Restoring a Model](#)
- white box models, [Making Predictions](#)
- Wide & Deep neural network, [Building Complex Models Using the Functional API](#)-[Building Complex Models Using the Functional API](#)
- window length, [Creating the Training Dataset](#)
- wisdom of the crowd, [Ensemble Learning and Random Forests](#)
- word embeddings, [Encoding Categorical Features Using Embeddings](#)
 - neural machine translation, [An Encoder–Decoder Network for Neural Machine Translation](#)-[Beam Search](#)
 - sentiment analysis, [Sentiment Analysis](#)-[Reusing Pretrained Embeddings and Language Models](#)
- worker task type, [Training a Model on a TensorFlow Cluster](#)
- workers, [Data parallelism with centralized parameters](#)

X

- Xavier initialization, [Glorot and He Initialization](#)
- Xception (Extreme Inception), [Xception-Xception](#), [Pretrained Models for Transfer Learning](#)-[Pretrained Models for Transfer Learning](#)
- XLA (accelerated linear algebra), [TensorFlow Functions and Graphs](#)
- XOR (exclusive or) problem, [The Perceptron](#)

Y

- You Only Look Once (YOLO), [You Only Look Once](#)-[You Only Look Once](#)

Z

- zero padding, [Convolutional Layers](#), [Implementing Convolutional Layers with Keras](#)
- zero-shot learning (ZSL), [An Avalanche of Transformer Models](#), [Vision Transformers](#)
- ZFNet, [AlexNet](#)

[Support](#) [Sign Out](#)

©2022 O'REILLY MEDIA, INC. [TERMS OF SERVICE](#) [PRIVACY POLICY](#)